arXiv:2611.04368v1 [cs.RO] 3 Nov 2026

RAFT: Reasoning-Aligned Fine-Tuning Improves Robustness of VLA Post-Training

Shresth Grover[1], **Andrey Kolobov**[2], **Dhruv Shah**[1]

[1]Princeton University, [2]Microsoft Research

Adapting a generalist vision-language-action (VLA) policy to a target task with limited demonstrations can improve nominal task success while reducing robustness to variations in appearance, objects, and instructions. We introduce reasoning-aligned fine-tuning (RAFT), a two-stage post-training procedure: RAFT first trains low-rank adapters on task-relevant reasoning annotations, then freezes them while fine-tuning the policy on action demonstrations. On a physical DROID robot, RAFT achieves 75.0% out-of-distribution success, compared with 27.9% for plain fine-tuning; on a multi-stage bimanual YAM task, it reaches 70.0% mean success, compared with 36.0% for the strongest baseline. RAFT also improves success under visual, object, and instruction shifts in simulated pick-and-place and open-drawer tasks. Compared with plain fine-tuning, RAFT policies place more attention on annotated task regions and yield higher linear-probe accuracy for object identity in later layers. A broader reasoning-data mixture further improves downstream success across demonstration budgets.

Keywords: Vision-language-action models, post-training, reasoning supervision

Website: princeton-prism.github.io/RAFT

Code: github.com/princeton-prism/RAFT

PRINCETON UNIVERSITY

1 Introduction

Generalist vision-language-action (VLA) policies, built on pretrained vision-language models and large robot datasets [1, 2, 4, 7, 22, 23], typically still require adaptation to reach the success rates needed to deploy the policy on a real-world target task [3, 24]. Because collecting task-specific demonstrations is expensive, this post-training often relies on a small dataset collected in a narrow set of environments. Fine-tuning can improve nominal success while reducing robustness to visual and semantic variations absent from those demonstrations [12, 21, 35, 37, 38, 43]. Existing approaches address this problem through parameter merging, as in RETAIN [39], or regularization during adaptation [17, 20]. We study a complementary approach, aligning a pretrained policy with task-relevant reasoning supervision before adapting it to the target actions.

Reasoning supervision can teach policies to identify task-relevant objects, spatial relations, and motion structure [10, 26, 41, 42]. We study how reasoning capabilities in pretrained VLMs can support low-data adaptation of generalist VLAs to a target skill. We introduce reasoning-aligned fine-tuning (RAFT), a two-stage adaptation procedure. Stage I learns low-rank adapters in the language and vision modules from task-relevant reasoning annotations while keeping pretrained parameters fixed. Stage II freezes these adapters and fine-tunes the language model, action expert, and projection layers on action demonstrations. Action gradients therefore do not directly update the reasoning-adapter weights. We find that how reasoning supervision is incorporated matters: jointly co-training reasoning and actions underperforms plain fine-tuning (63.2% vs. 79.6% mean success on simulated pick and

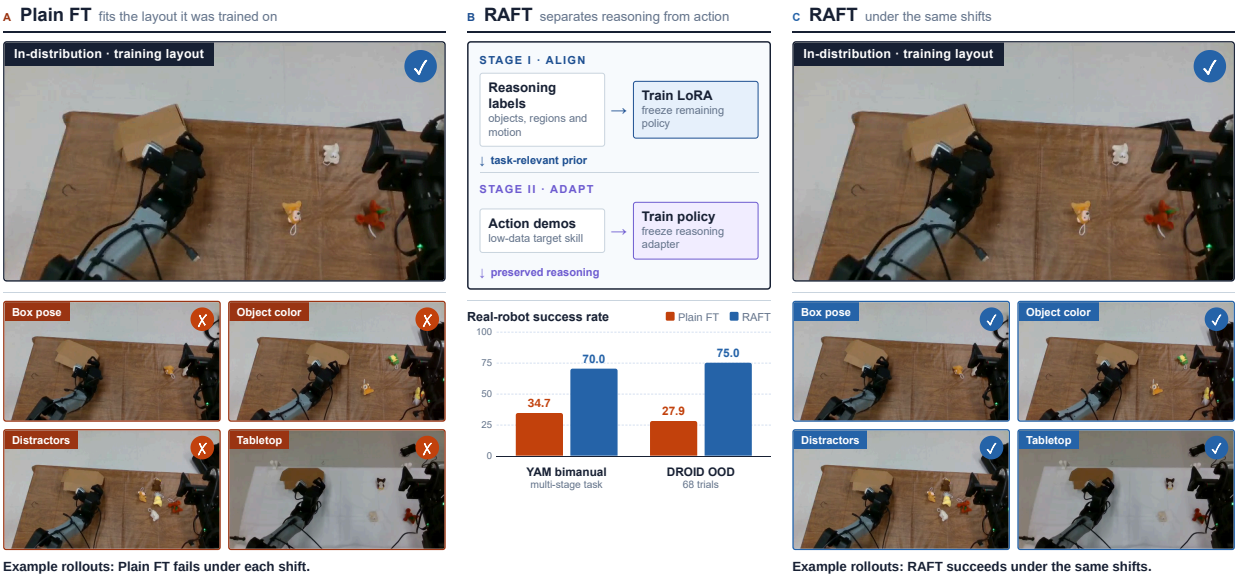

Figure 1 | Reasoning-aligned fine-tuning for robust VLA adaptation. (a) Plain FT, fine-tuned only on actions, succeeds in the training layout but fails under small scene changes. (b) RAFT first learns task-relevant reasoning through LoRA adapters, then freezes them while fine-tuning the policy on action demonstrations. (c) With the same demonstrations, RAFT succeeds under these shifts. On physical robots, RAFT raises success over Plain FT from 34.7% to 70.0% on YAM and from 27.9% to 75.0% on out-of-distribution DROID trials.

place), and updating the reasoning adapter during action learning falls below Plain FT (74.0% vs. 79.6%) (Secs. C.3 and C.2). Staging the two, with the adapter frozen, yields these robustness gains. RAFT policies also place more attention on annotated task regions and yield higher linear-probe accuracy for object identity and episode-time position than plainly fine-tuned policies (Sec. 4.3).

We evaluate RAFT on two simulated tasks and two physical platforms: the single-arm DROID setup [22] and the bimanual YAM setup [10]. Evaluations test visual, object, instruction, and geometric variations of the adapted tasks. RAFT achieves the highest mean success on simulated pick and place and improves open-drawer success, with the largest gains on the physical robots: success on 75% of out-of-distribution DROID trials versus 28% for plain fine-tuning, and 70% mean success on the multi-stage YAM task versus 36% for the strongest baseline. Our contributions are threefold: (1) We propose RAFT a post training recipt that separates task relevant reasoning supervision from action fine tuning (2) We show improved downstream performace and robustness in simulation and on two physical robot platforms; and (3) We analyze why RAFT helps generalist policy adapt to a target by analyzing the features learned by RAFT.

2 Related Work

Generalist robot policies have made substantial progress in learning transferable skills and supporting deployment through post-training [5, 15, 23, 27, 30, 34, 40]. Within this setting, work on VLA fine-tuning studies reliable task execution [24, 43], while FAST improves action tokenization for autoregressive VLAs [31]. ControlVLA uses object-centric conditioning for few-shot adaptation [28]. These advances make specialization more effective, but nominal task success can still mask failures under distribution shift [11, 12, 44]. This motivates studying robustness to variations absent from

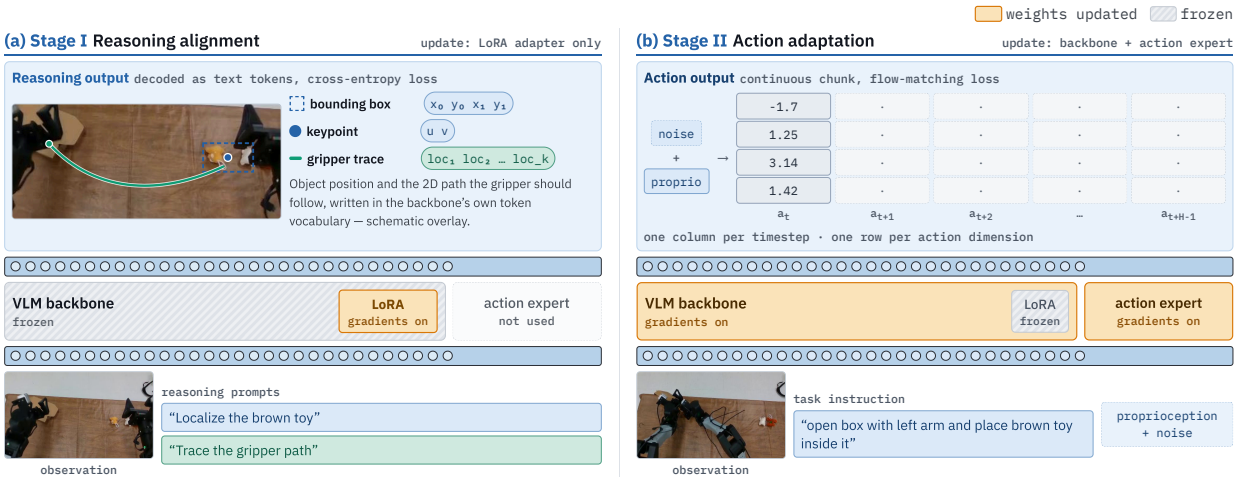

Figure 2 | Overview of RAFT. Stage I updates only the reasoning adapters, using reasoning annotations. Stage II freezes the adapters and updates the language model, action expert, and projection layers on action demonstrations, so action gradients do not reach the reasoning adapters.

the target demonstrations.

Robust post-training methods seek to adapt generalist policies while retaining useful pretrained capabilities. Weight-space approaches interpolate or merge pretrained and adapted parameters [9, 18, 36], while regularization and steerability-preserving adaptation constrain changes during fine-tuning [14, 17, 20, 33]. RETAIN [39] interpolates pretrained and fine-tuned policy weights to improve generalization to variations of a new skill and retain broader capabilities. DeLock [17] regularizes visual-encoder weight drift during post-training and uses contrastive prompt guidance at inference. RAFT instead introduces task-relevant reasoning alignment before action fine-tuning, with complementary parameter freezing across the two stages.

Reasoning and grounding supervision provide a foundation for this approach. Embodied Chain-of-Thought interleaves sub-goal, spatial, and motion reasoning with action prediction [41]. Related work incorporates spatial representations, affordances, keypoints, and manipulable parts [10, 26, 42], or object- and language-grounded supervision [4]. Steerable Policies use commands at multiple abstraction levels to control [6], while VLS uses VLM-generated rewards to steer frozen policies at inference [29]. Reasoning supervision is used to prepare that policy for specialization, and the learned adapters are held fixed during subsequent action learning. To the best of our knowledge, reasoning supervision has not been studied for low-data post-training of a pretrained generalist VLA. We evaluate this adaptation schedule through downstream robustness, additional training controls, and analyses of task-relevant representations.

3 Method

3.1 Overview and Problem Setup

We consider task adaptation from a pretrained generalist policy using a small demonstration dataset $\mathcal{D}_\eta = \{(s_t, a_t, T)\}$ collected in a narrow set of environments. Here s_t contains the policy observations, a_t is the demonstrated action, and T is the task instruction. We denote the Stage-II trainable non-adapter parameters (the language model, action expert, and projection layers) by θ, initialized from

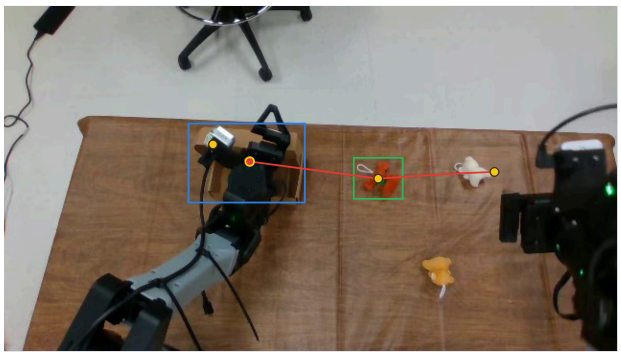

Figure 3 | Example Stage-I reasoning annotation. In this YAM example, the left arm holds the lid open while the right arm reaches for the toy. The target combines language, object boxes, keypoints, motion traces, and per-arm state information. Blue and green mark the box and toy, yellow marks keypoints, and red marks the active arm's trace.

θ_{pre}, and the LoRA parameters by ϕ [16]. In RAFT, the pretrained SigLIP weights remain frozen throughout adaptation, while LoRA modules in both the vision encoder and language model are trained in Stage I. Stage II freezes these learned adapters and updates θ. Thus, vision adaptation in RAFT occurs through Stage-I LoRA rather than updates to the pretrained SigLIP weights. The resulting action policy is $\pi_{\theta,\phi}(a_t \mid s_t, T)$. Figure 2 summarizes the two stages; Sec. A gives the adapter placement and optimization settings.

3.2 Stage I: Reasoning Alignment

Stage I trains the adapter to predict task-relevant reasoning annotations. A VLM annotator (Gemini 3.1 Pro) produces task-relevant reasoning targets from observations and instructions on $\mathcal{D}_\eta$, yielding $\mathcal{D}_{\text{reason}} = \{(s, r, T)\}$. These annotations describe the objects and interactions relevant to the task. The annotation choices include bounding boxes, keypoints, traces, spatial descriptions, and visual question answering. Section C.1 compares their effects. Each annotation is represented as a target token sequence $r = (r_1, \ldots, r_{|r|})$. We train the LoRA adapter using standard autoregressive cross-entropy:

$$\mathcal{L}_{\text{reason}}(\phi) = -\mathbb{E}_{(s,r,T)\sim\mathcal{D}_{\text{reason}}} \sum_{j=1}^{|r|} \log p_{\theta_{\text{pre}},\phi}(r_j \mid s, T, r_{<j}). \tag{1}$$

Only ϕ is updated in this stage; all remaining parameters are frozen. This supervision encourages task-relevant grounding before action adaptation.

3.3 Stage II: Action Adaptation

Stage II turns the reasoning-aligned initialization into an action policy for the target task. We retain the learned adapter $\phi^\star$ and freeze its parameters. The Stage-II parameter set θ is then updated by

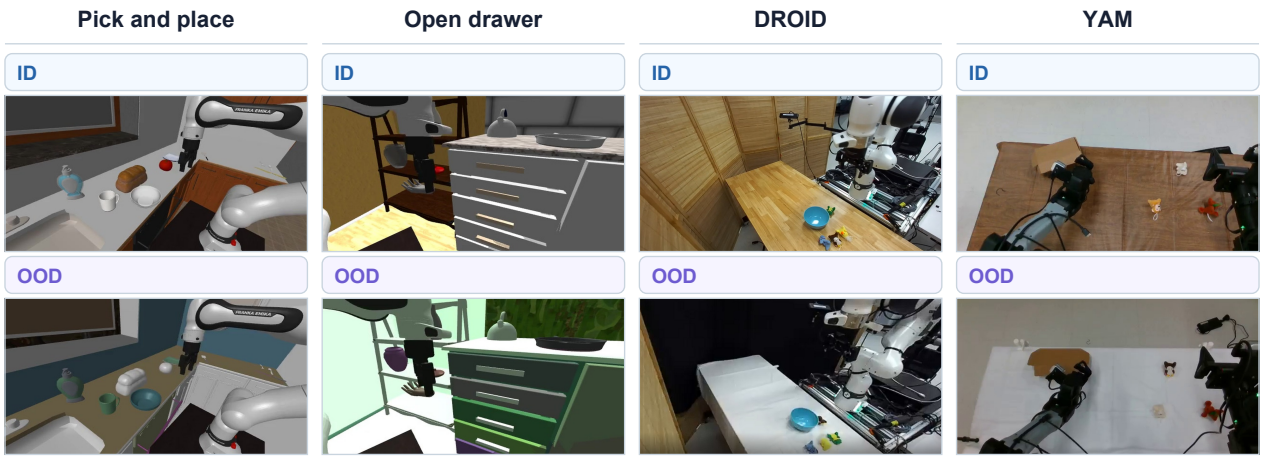

Figure 4 | Evaluation scenes. Representative in-distribution (ID) and out-of-distribution (OOD) observations, left to right: simulated pick and place, simulated open drawer, physical DROID pick and place, and the physical YAM two-stage task. The OOD panels show texture changes in simulation and background changes on the physical robots. The full set of evaluated shifts is shown in Fig. 20.

supervised fine-tuning (SFT) on the task demonstrations [23, 24]:

$$\theta^\star = \arg\min_\theta \underbrace{\mathbb{E}_{(s,a,T)\sim\mathcal{D}_\eta}\left[\ell_{\text{SFT}}(\pi_{\theta,\phi^\star}; s, a, T)\right]}_{\mathcal{L}_{\text{BC}}(\theta;\mathcal{D}_\eta,\phi^\star)}. \tag{2}$$

Here ℓ_{SFT} is the underlying policy's action-supervision objective. Stage II updates the language model, action expert, and projection layers while the adapter and SigLIP vision tower remain frozen. At deployment, the resulting policy is $\pi_{\theta^\star,\phi^\star}$. The adapter ablation in the appendix (Sec. C.2) evaluates the benefit of freezing the reasoning adapter during action learning. Attention and linear-probe analyses (Sec. 4.3) provide empirical evidence of retained task grounding in the resulting policy.

4 Experiments

The goal of our experiments is to evaluate RAFT's ability to retain task-relevant information and improve downstream robustness. Concretely, we aim to answer: (1) does RAFT generalize better than baselines under distribution shift, in simulation and on physical robots? (2) do RAFT's representations retain more task-relevant information than a plainly fine-tuned policy, and does this track downstream success? (3) how does downstream performance scale with the amount of reasoning data used in Stage I? and (4) How freezing Stage-I adapted compares with common approaches like co-training simulatenously?

4.1 Experimental setup

Baselines We compare RAFT with standard supervised fine-tuning (*Plain FT*), *RETAIN* [39], which merges pretrained and fine-tuned weights, and *DeLock* [17], which regularizes visual-encoder weight drift during fine-tuning and applies contrastive prompt guidance at inference. We additionally report *Base*, the pretrained generalist policy evaluated without post-training. All fine-tuned policies are evaluated after 5,000 action-training steps. Sec. A.3 gives the baseline settings; Sec. 4.5 ablates the training design, with full adapter and co-training results in the appendix (Secs. C.2 and C.3).

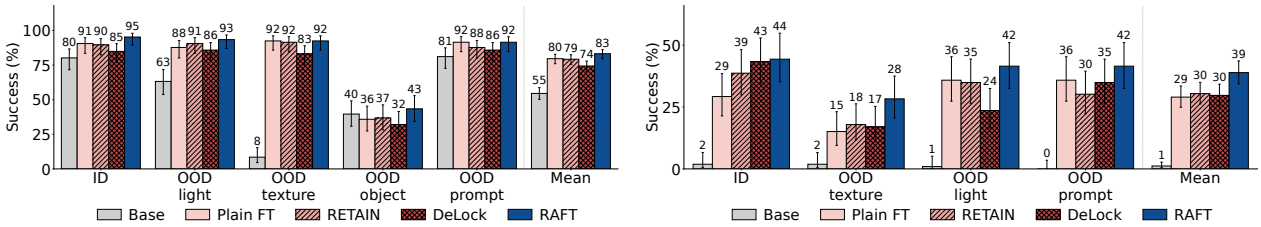

Figure 5 | Simulation results under distribution shift. Success rates for pick and place (left) and open drawer (right), with 106 trials per split. RAFT uses full reasoning for pick and place and bounding-box supervision for open drawer (Sec. C.1). RAFT has the highest mean success on both tasks: 83% versus 80% for Plain FT on pick and place, and 39% versus 29% on open drawer.

4.2 RAFT generalizes under distribution shift

We evaluate RAFT on pick-and-place and open-drawer tasks in simulation and on two real-world platforms: the single-arm DROID setup [22] and the dexterous, dual-arm YAM setup [10], under the visual, object, and instruction or geometric shifts described below.

Simulation We evaluate MolmoBot [8] in the MolmoSpaces simulator [25] on pick-and-place and open-drawer tasks. For each task, we use 106 successful demonstrations from one layout provided in the MolmoSpaces benchmark for post-training. We measure robustness using perturbations absent from the demonstrations. *OOD-light* changes the number, position, and intensity of light sources; *OOD-texture* replaces table and scene textures; *OOD-object* substitutes a semantically similar manipulation object unseen during post-training and is evaluated only for pick and place; and *OOD-prompt* paraphrases the instruction without changing its meaning. *ID* uses the demonstration layout. Each evaluation split contains $n = 106$ trials with matched seeds across methods.

RAFT has the highest simulated success, with the largest margin on novel objects On pick and place, RAFT achieves 83% mean success across the five evaluation splits, compared with 80% for Plain FT, 79% for RETAIN, and 74% for DeLock (Fig. 5). The largest margin is under object shift: 43% for RAFT versus 36% for Plain FT and 37% for RETAIN. On the open drawer task, RAFT reaches 44% success in ID and 28% under texture shift, versus 29% and 15% for Plain FT, and it also leads under lighting and instruction shifts.

Real-world: DROID We post-train MolmoBot [8] on 50 pick-and-place demonstrations and evaluate on a physical DROID robot under four axes: the demonstration setup (ID; blue, yellow, and brown toys), novel object colors (OOD-Object, e.g., a white toy in place of the brown one), added distractors (OOD-Distractor), and a changed tabletop surface (OOD-Background). On out-of-distribution DROID trials, RAFT succeeds in 75.0% of trials versus 27.9% for Plain FT (Fig. 6): 83% versus 33% with novel object colors, 75% versus 25% with distractors, and 65% versus 25% on the changed tabletop. In the demonstration setup, the two are close (63% versus 60%).

Real-world: YAM For the dual-arm YAM platform, we post-train $\pi_{0.5}$ [4] on 100 demonstrations of a long-horizon task: open a hinged box, pick up an object, and place it inside. Because placement depends on the box being open, the task separates completing the first interaction from completing the whole sequence. We evaluate in the demonstration setup and under four shifts: clutter, box pose, object color, and background. On the bimanual YAM task, RAFT completes the full sequence in 70.0% of trials pooled over the five axes, versus 36% for DeLock, 35% for Plain FT, and 29% for RETAIN

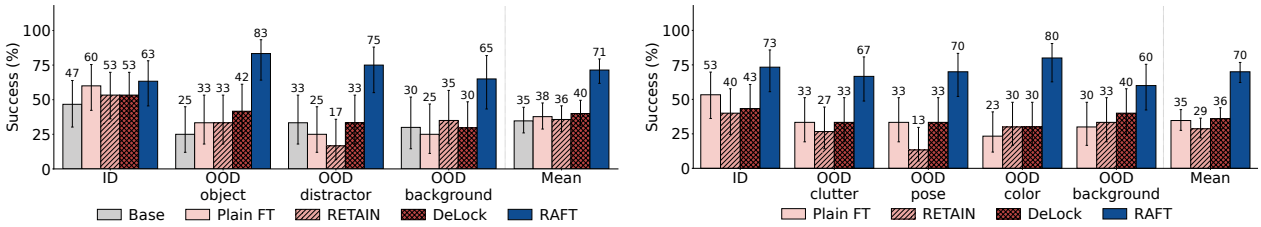

Figure 6 | Real-robot results under distribution shift. Left: DROID pick and place. Right: YAM box opening and place inside box. RAFT has the highest success on both platforms: 75% versus 28% for Plain FT on out-of-distribution DROID trials, and 70% versus 35% on YAM.

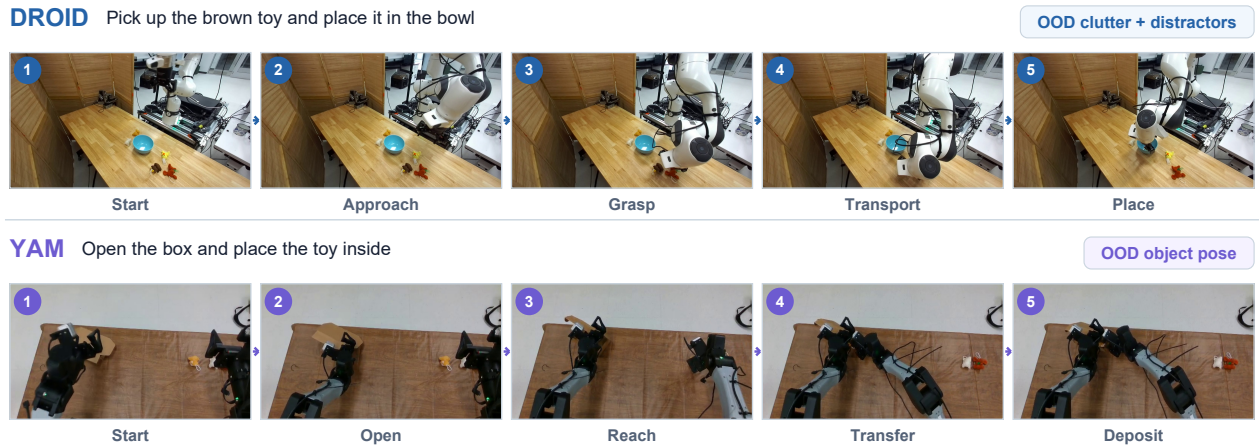

Figure 7 | Successful physical-robot rollouts. Ordered frames show DROID picking up the brown toy and placing it in the bowl (top) and YAM opening the box and placing the toy inside (bottom). Numbers indicate temporal order; circles and arrows mark the manipulated object and motion.

(Fig. 6), and it has the highest success on every axis (73% in-distribution and 60%–80% under the four shifts). Plain FT usually opens the box but then fails to place the object.

4.3 Representation Analysis

We examine representation dimensionality, linearly decodable task information, and visual attention to assess what RAFT preserves during post-training. The measurement protocol is in Sec. E.

Late-layer dimensionality correlates with downstream success We measure the effective dimensionality of the backbone's hidden states with the participation ratio (PR) [13, 19, 32]:

$$\mathrm{PR} = \frac{(\sum_i \lambda_i)^2}{\sum_i \lambda_i^2}, \tag{3}$$

where λ_i are the eigenvalues of the covariance of mean-pooled hidden states across frames. Across six Stage-I reasoning variants, late-layer PR correlates with success on each task-matched evaluation axis; at layer 15, absolute PR has mean Spearman $\rho = 0.76$ (Fig. 19).

Object identity and task progress are more linearly decodable from RAFT's late layers We fit linear probes on the same late layers to predict object identity (seven classes, chance 1/7) and episode-time quartile (four bins, chance 1/4). In layers 14–17, probes on RAFT reach 0.51–0.53 balanced accuracy

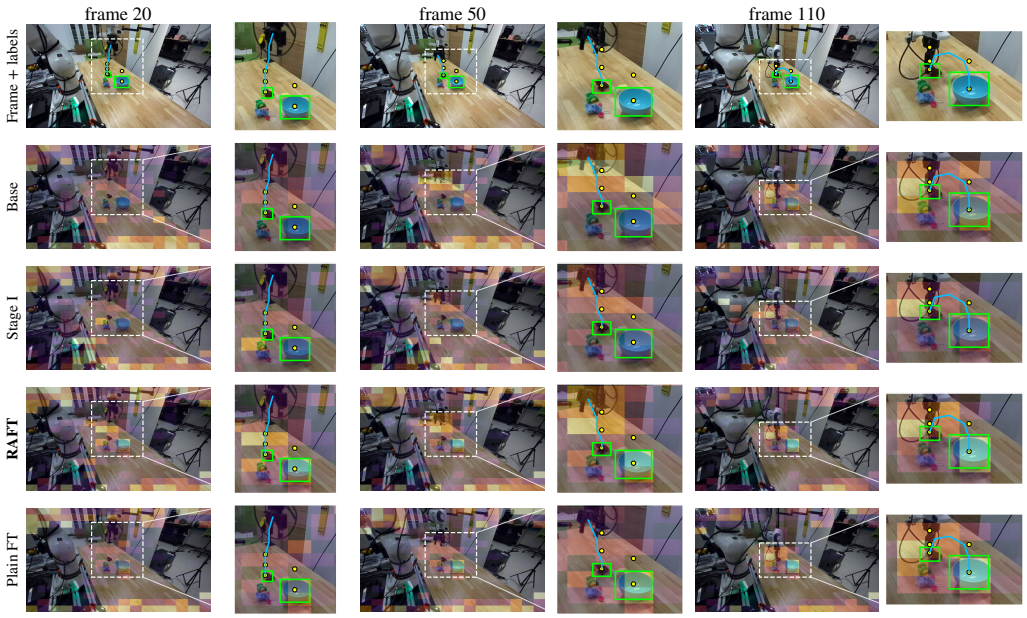

Figure 8 | Attention maps along one trajectory. Columns show successive observations; rows compare Base, Stage I, RAFT, and Plain FT. Overlays identify annotated task regions, and color scales are shared within each column. RAFT places more attention on the annotated regions than Plain FT; Fig. 10 quantifies this across layers.

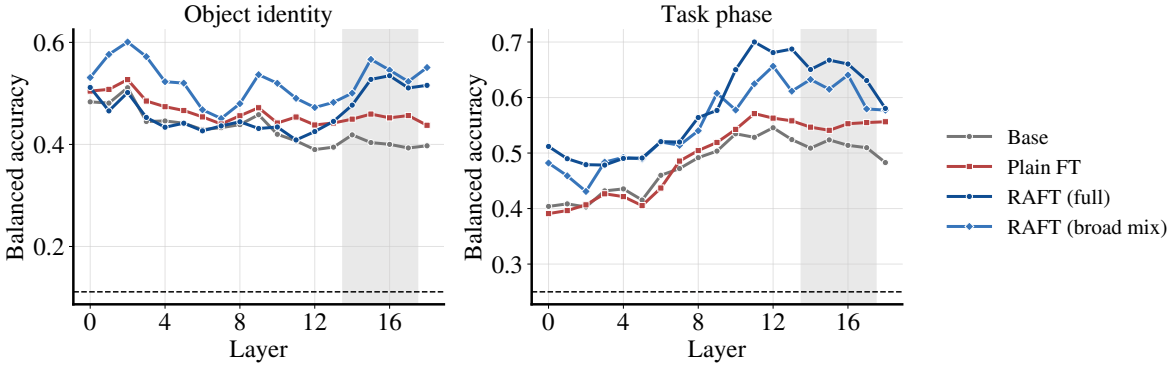

Figure 9 | Layer-wise linear probes. Balanced accuracy for object identity (left; chance 1/7) and episode-time quartile (right; chance 1/4). Policies use the same 414 frames; episodes are held out from probe fitting, not policy training. Shading marks layers 14–17. In these layers, RAFT reaches 0.51–0.53 object-identity accuracy versus 0.45 for Plain FT, and 0.62–0.65 quartile accuracy versus 0.55.

for object identity, against 0.45 for Plain FT and 0.40 for the base policy, and 0.62–0.65 for episode quartile, against 0.55 and 0.51 (Fig. 9). Object-identity decoding also rises sharply between layers 13 and 15 for RAFT, but barely changes for Plain FT or the base policy. The layers where PR tracks success are the same layers where this information becomes decodable.

RAFT has higher attention enrichment on annotated task regions We measure how much of the policy's visual attention falls on annotated task regions, such as the manipulated object, the target, and the gripper, relative to their share of the image (enrichment; values above one mean attention is concentrated there). At layer 11, RAFT gives these regions 3.2–3.4 times their uniform share of attention, against 1.9–2.1 for Plain FT, while attention to the background stays near uniform (Figs. 8

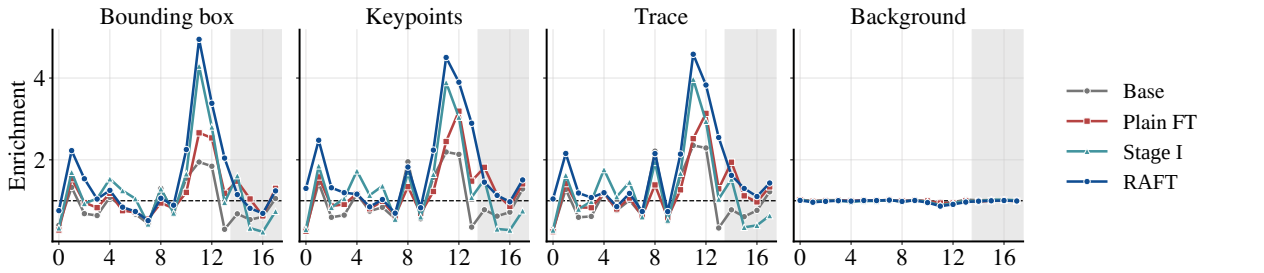

Figure 10 | Layer-wise attention enrichment. Enrichment for object boxes, keypoints, traces, and background; values above one indicate attention exceeding a region's image-area share. Curves compare Base, Plain FT, Stage I, and RAFT; shading marks layers 14–17. At layer 11, RAFT gives task regions 3.2–3.4 times their area share, versus 1.9–2.1 for Plain FT, while background attention stays near uniform.

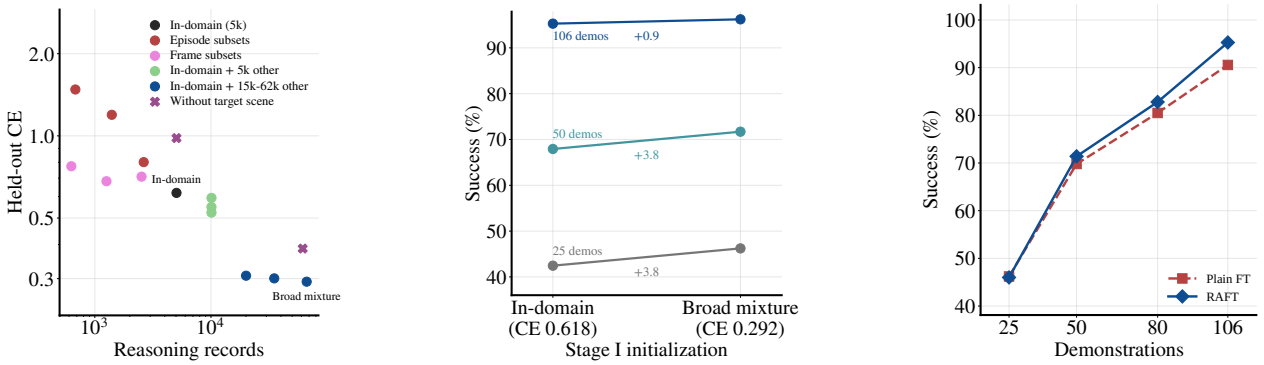

Figure 11 | Effect of reasoning data and demonstration budget. Left: held-out reasoning loss across corpus sizes and constructions. Middle: in-distribution pick-and-place success after 5,000 Stage-II steps from in-domain and broader-mixture Stage-I initializations. Right: RAFT versus Plain FT across demonstration budgets, with 106 evaluation episodes per budget. The broader mixture lowers held-out loss (0.29 versus 0.62) and improves this success by 0.9–3.8 points at every budget.

and 10); the separation narrows in later layers. The concentration persists through Stage-II action training and stays above that of plain fine-tuning. The size of these effects depends on the reasoning modality used in Stage I; Sec. C.1 reports the full modality ablation.

4.4 Effect of reasoning data on held-out loss and downstream success

More reasoning data lowers held-out loss On a common 560-record held-out split, cross-entropy on full-reasoning targets falls as the number of reasoning records grows, in both frame- and episode-subsampled series (Fig. 11, left). Corpus composition also matters: 628 frame-subsampled records reach a loss of 0.77, while 680 episode-subsampled records reach only 1.48. The broader mixture, which adds roughly 60,000 records to the 5,000-record in-domain corpus, gives the lowest loss (0.29, versus 0.62 in-domain) and still reaches 0.39 with target-simulator records removed.

Broader reasoning data also helps downstream Initializing Stage II from the broader-mixture adapter instead of the in-domain one improves in-distribution success by 0.9–3.8 points at each of the three demonstration budgets (Fig. 11, middle). Against Plain FT at matched budgets (right panel), the two methods tie with 25 demonstrations (46.2% success each), and RAFT leads modestly at intermediate budgets. With all demonstrations, RAFT reaches 95.3% ID success against 90.6% for Plain FT (Fig. 5).

4.5 Effect of unfrozen adapter

We ablate two design choices of RAFT on simulated pick and place: how reasoning and action supervision are combined, and whether the reasoning adapter is updated during action learning (Fig. 12). *Co-train* trains on reasoning and action examples jointly, sampling them equally and weighting their losses 1:1. *Adapter trainable* runs both stages but keeps the Stage-I adapter trainable in Stage II, and *Stage I only* applies the adapter without action fine-tuning.

Across five matched splits with 106 episodes each, co-training reaches 63.21% mean success, below both Plain FT (79.62%) and RAFT (83.21%). Updating the adapter in Stage II lowers success to 73.96%, and Stage I alone reaches 48.30%. Plain FT with a frozen vision encoder also stays at 79.62%, so freezing the pretrained vision weights alone does not account for the gain. Split-level results are in Tables 3 and 4.

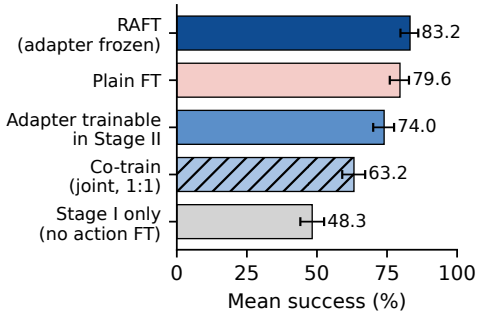

Figure 12 | Training-scheme ablation on simulated pick and place. Mean success over trials. Staged training with a frozen adapter achieves the highest success (83.2%), above Plain FT (79.6%), a trainable adapter (74.0%), and co-training (63.2%).

5 Conclusion

We presented RAFT, which first trains a low-rank adapter on task-relevant reasoning annotations, freezes it, and only then fine-tunes the pretrained VLA on actions. Across two simulated tasks and two real robots, this staging improved success under distribution shift relative to plain fine-tuning, RETAIN, and DeLock. We also show how training reasoning adapter on different modalities can lead to varied performance gains and selection of optimal reasoning modality for a specific task remains an open question. Since, RAFT approaches post-training from a different perspective to current approaches, it makes RAFT compatible with these approaches and future work could explore utlizing a combination of RAFT and existing methods to do more robust post-training.

Acknowledgments

This research was partially supported by the Toyota Research Institute and Microsoft Research, with compute support from Google TPU Research Cloud, NVIDIA Academic Grant Program, and Gemini Academic Program. The authors also thank Samuel M. Bateman and Lihan Zha for helpful discussions.

References

[1] Rohan Anil, Andrew M Dai, Orhan Firat, Melvin Johnson, Dmitry Lepikhin, Alexandre Passos, Siamak Shakeri, Emanuel Taropa, Paige Bailey, Zhifeng Chen, et al. PaLM 2 technical report. *arXiv preprint arXiv:2305.10403*, 2023.

[2] Johan Bjorck, Fernando Castañeda, Nikita Cherniadev, Xingye Da, Runyu Ding, Linxi Fan, Yu Fang, Dieter Fox, Fengyuan Hu, Spencer Huang, et al. GR00T N1: An open foundation model for generalist humanoid robots. *arXiv preprint arXiv:2503.14734*, 2025.

[3] Kevin Black, Noah Brown, Danny Driess, Adnan Esmail, Michael Equi, Chelsea Finn, Niccolo Fusai, Lachy Groom, Karol Hausman, Brian Ichter, et al. π_0: A vision-language-action flow model for general robot control. *arXiv preprint arXiv:2410.24164*, 2024.

[4] Kevin Black, Noah Brown, James Darpinian, Karan Dhabalia, Danny Driess, Adnan Esmail, Michael Robert Equi, Chelsea Finn, Niccolo Fusai, Manuel Y. Galliker, Dibya Ghosh, Lachy Groom, Karol Hausman, Brian Ichter, Szymon Jakubczak, Tim Jones, Liyiming Ke, Devin LeBlanc, Sergey Levine, Adrian Li-Bell, Mohith Mothukuri, Suraj Nair, Karl Pertsch, Allen Z. Ren, Lucy Xiaoyang Shi, Laura Smith, Jost Tobias Springenberg, Kyle Stachowicz, James Tanner, Quan Vuong, Homer Walke, Anna Walling, Haohuan Wang, Lili Yu, and Ury Zhilinsky. $\pi_{0.5}$: a vision-language-action model with open-world generalization. In Joseph Lim, Shuran Song, and Hae-Won Park, editors, *Proceedings of The 9th Conference on Robot Learning*, volume 305 of *Proceedings of Machine Learning Research*, pages 17–40. PMLR, 27–30 Sep 2025. URL https://proceedings.mlr.press/v305/black25a.html.

[5] Konstantinos Bousmalis, Giulia Vezzani, Dushyant Rao, Coline Devin, Alex X Lee, Maria Bauzá, Todor Davchev, Yuxiang Zhou, Agrim Gupta, Akhil Raju, et al. Robocat: A self-improving generalist agent for robotic manipulation. *arXiv preprint arXiv:2306.11706*, 2023.

[6] William Chen, Jagdeep Singh Bhatia, Catherine Glossop, Nikhil Mathihalli, Ria Doshi, Andy Tang, Danny Driess, Karl Pertsch, and Sergey Levine. Steerable vision-language-action policies for embodied reasoning and hierarchical control. *arXiv preprint arXiv:2602.13193*, 2026.

[7] Open X-Embodiment Collaboration, Abby O'Neill, Abdul Rehman, Abhinav Gupta, Abhiram Maddukuri, Abhishek Gupta, Abhishek Padalkar, Abraham Lee, Acorn Pooley, Agrim Gupta, Ajay Mandlekar, Ajinkya Jain, Albert Tung, Alex Bewley, Alex Herzog, Alex Irpan, Alexander Khazatsky, Anant Rai, Anchit Gupta, Andrew Wang, Andrey Kolobov, Anikait Singh, Animesh Garg, Aniruddha Kembhavi, Annie Xie, Anthony Brohan, Antonin Raffin, Archit Sharma, Arefeh Yavary, Arhan Jain, Ashwin Balakrishna, Ayzaan Wahid, Ben Burgess-Limerick, Beomjoon Kim, Bernhard Schölkopf, Blake Wulfe, Brian Ichter, Cewu Lu, Charles Xu, Charlotte Le, Chelsea Finn, Chen Wang, Chenfeng Xu, Cheng Chi, Chenguang Huang, Christine Chan, Christopher Agia, Chuer Pan, Chuyuan Fu, Coline Devin, Danfei Xu, Daniel Morton, Danny Driess, Daphne Chen, Deepak Pathak, Dhruv Shah, Dieter Büchler, Dinesh Jayaraman, Dmitry Kalashnikov, Dorsa Sadigh, Edward Johns, Ethan Foster, Fangchen Liu, Federico Ceola, Fei Xia, Feiyu Zhao, Felipe Vieira Frujeri, Freek Stulp, Gaoyue Zhou, Gaurav S. Sukhatme, Gautam Salhotra, Ge Yan, Gilbert Feng, Giulio Schiavi, Glen Berseth, Gregory Kahn, Guangwen Yang, Guanzhi Wang, Hao Su, Hao-Shu Fang, Haochen Shi, Henghui Bao, Heni Ben Amor, Henrik I Christensen, Hiroki Furuta, Homanga Bharadhwaj, Homer Walke, Hongjie Fang, Huy Ha, Igor Mordatch, Ilija Radosavovic, Isabel Leal, Jacky Liang, Jad Abou-Chakra, Jaehyung Kim, Jaimyn Drake, Jan Peters, Jan Schneider, Jasmine Hsu, Jay Vakil, Jeannette Bohg, Jeffrey Bingham, Jeffrey Wu, Jensen Gao, Jiaheng Hu, Jiajun Wu, Jialin Wu, Jiankai Sun, Jianlan Luo, Jiayuan Gu, Jie Tan, Jihoon Oh, Jimmy Wu, Jingpei Lu, Jingyun Yang, Jitendra Malik, João Silvério, Joey Hejna, Jonathan Booher, Jonathan Tompson, Jonathan Yang, Jordi Salvador, Joseph J. Lim, Junhyek Han, Kaiyuan Wang, Kanishka Rao, Karl Pertsch, Karol Hausman, Keegan Go, Keerthana Gopalakrishnan, Ken Goldberg, Kendra Byrne, Kenneth Oslund, Kento Kawaharazuka, Kevin Black, Kevin Lin, Kevin Zhang, Kiana Ehsani, Kiran Lekkala, Kirsty Ellis, Krishan Rana, Krishnan Srinivasan, Kuan Fang, Kunal Pratap Singh, Kuo-Hao Zeng, Kyle Hatch, Kyle Hsu, Laurent Itti, Lawrence Yunliang Chen, Lerrel Pinto, Li Fei-Fei, Liam Tan, Linxi "Jim" Fan, Lionel Ott, Lisa Lee, Luca Weihs, Magnum Chen, Marion Lepert, Marius Memmel, Masayoshi Tomizuka, Masha Itkina, Mateo Guaman Castro, Max Spero, Maximilian Du, Michael Ahn, Michael C. Yip, Mingtong Zhang, Mingyu Ding, Minho Heo, Mohan Kumar Srirama, Mohit Sharma, Moo Jin Kim, Muhammad Zubair

Irshad, Naoaki Kanazawa, Nicklas Hansen, Nicolas Heess, Nikhil J Joshi, Niko Suenderhauf, Ning Liu, Norman Di Palo, Nur Muhammad Mahi Shafiullah, Oier Mees, Oliver Kroemer, Osbert Bastani, Pannag R Sanketi, Patrick "Tree" Miller, Patrick Yin, Paul Wohlhart, Peng Xu, Peter David Fagan, Peter Mitrano, Pierre Sermanet, Pieter Abbeel, Priya Sundaresan, Qiuyu Chen, Quan Vuong, Rafael Rafailov, Ran Tian, Ria Doshi, Roberto Martín-Martín, Rohan Baijal, Rosario Scalise, Rose Hendrix, Roy Lin, Runjia Qian, Ruohan Zhang, Russell Mendonca, Rutav Shah, Ryan Hoque, Ryan Julian, Samuel Bustamante, Sean Kirmani, Sergey Levine, Shan Lin, Sherry Moore, Shikhar Bahl, Shivin Dass, Shubham Sonawani, Shubham Tulsiani, Shuran Song, Sichun Xu, Siddhant Haldar, Siddharth Karamcheti, Simeon Adebola, Simon Guist, Soroush Nasiriany, Stefan Schaal, Stefan Welker, Stephen Tian, Subramanian Ramamoorthy, Sudeep Dasari, Suneel Belkhale, Sungjae Park, Suraj Nair, Suvir Mirchandani, Takayuki Osa, Tanmay Gupta, Tatsuya Harada, Tatsuya Matsushima, Ted Xiao, Thomas Kollar, Tianhe Yu, Tianli Ding, Todor Davchev, Tony Z. Zhao, Travis Armstrong, Trevor Darrell, Trinity Chung, Vidhi Jain, Vikash Kumar, Vincent Vanhoucke, Vitor Guizilini, Wei Zhan, Wenxuan Zhou, Wolfram Burgard, Xi Chen, Xiangyu Chen, Xiaolong Wang, Xinghao Zhu, Xinyang Geng, Xiyuan Liu, Xu Liangwei, Xuanlin Li, Yansong Pang, Yao Lu, Yecheng Jason Ma, Yejin Kim, Yevgen Chebotar, Yifan Zhou, Yifeng Zhu, Yilin Wu, Ying Xu, Yixuan Wang, Yonatan Bisk, Yongqiang Dou, Yoonyoung Cho, Youngwoon Lee, Yuchen Cui, Yue Cao, Yueh-Hua Wu, Yujin Tang, Yuke Zhu, Yunchu Zhang, Yunfan Jiang, Yunshuang Li, Yunzhu Li, Yusuke Iwasawa, Yutaka Matsuo, Zehan Ma, Zhuo Xu, Zichen Jeff Cui, Zichen Zhang, Zipeng Fu, and Zipeng Lin. Open X-Embodiment: Robotic learning datasets and RT-X models. https://arxiv.org/abs/2310.08864, 2023.

[8] Abhay Deshpande, Maya Guru, Rose Hendrix, Snehal Jauhri, Ainaz Eftekhar, Rohun Tripathi, Max Argus, Jordi Salvador, Haoquan Fang, Matthew Wallingford, et al. MolmoB0T: Large-scale simulation enables zero-shot manipulation. *arXiv preprint arXiv:2603.16861*, 2026.

[9] Sombit Dey, Jan-Nico Zaech, Nikolay Nikolov, Luc Van Gool, and Danda Pani Paudel. ReVLA: Reverting visual domain limitation of robotic foundation models. In *IEEE International Conference on Robotics and Automation (ICRA)*, 2025.

[10] Haoquan Fang, Jiafei Duan, Donovan Clay, Sam Wang, Shuo Liu, Weikai Huang, Xiang Fan, Wei-Chuan Tsai, Shirui Chen, Yi Ru Wang, et al. MolmoAct2: Action reasoning models for real-world deployment. *arXiv preprint arXiv:2605.02881*, 2026.

[11] Senyu Fei, Siyin Wang, Junhao Shi, Zihao Dai, Jikun Cai, Pengfang Qian, Li Ji, Xinzhe He, Shiduo Zhang, Zhaoye Fei, et al. LIBERO-Plus: A progressive robustness benchmark for visual-language-action models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 38574–38583, 2026.

[12] Jensen Gao, Suneel Belkhale, Sudeep Dasari, Ashwin Balakrishna, Dhruv Shah, and Dorsa Sadigh. A taxonomy for evaluating generalist robot manipulation policies. *IEEE Robotics and Automation Letters*, 2026.

[13] Peiran Gao, Eric Trautmann, Byron Yu, Gopal Santhanam, Stephen Ryu, Krishna Shenoy, and Surya Ganguli. A theory of multineuronal dimensionality, dynamics and measurement. *bioRxiv*, 2017. doi: 10.1101/214262.

[14] Shresth Grover, Akshay Gopalkrishnan, Bo Ai, Henrik I Christensen, Hao Su, and Xuanlin Li. Enhancing generalization in vision-language-action models by preserving pretrained representations. *arXiv preprint arXiv:2509.11417*, 2025.

[15] Yanjiang Guo, Jianke Zhang, Xiaoyu Chen, Xiang Ji, Yen-Jen Wang, Yucheng Hu, and Jianyu Chen. Improving vision-language-action model with online reinforcement learning. In *2025 IEEE International Conference on Robotics and Automation (ICRA)*, pages 15665–15672. IEEE, 2025.

[16] Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations*, 2022. URL https://openreview.net/forum?id=nZeVKeeFYf9.

[17] Suning Huang, Jiaqi Shao, Ke Wang, Qianzhong Chen, Jiankai Sun, Yanjiang Guo, Mac Schwager, and

Jeannette Bohg. Breaking lock-in: Preserving steerability under low-data VLA post-training. *arXiv preprint arXiv:2604.23121*, 2026.

[18] Dong-Hwan Jang, Sangdoo Yun, and Dongyoon Han. Model stock: All we need is just a few fine-tuned models. In *European Conference on Computer Vision*, pages 207–223. Springer, 2024.

[19] Li Jing, Pascal Vincent, Yann LeCun, and Yuandong Tian. Understanding dimensional collapse in contrastive self-supervised learning. In *International Conference on Learning Representations (ICLR)*, 2022.

[20] Nikita Kachaev, Mikhail Kolosov, Daniil Zelezetsky, Alexey K Kovalev, and Aleksandr I Panov. Don't blind your VLA: Aligning visual representations for OOD generalization. *arXiv preprint arXiv:2510.25616*, 2025.

[21] Christos Kaplanis, Murray Shanahan, and Claudia Clopath. Policy consolidation for continual reinforcement learning. *arXiv preprint arXiv:1902.00255*, 2019.

[22] Alexander Khazatsky, Karl Pertsch, Suraj Nair, Ashwin Balakrishna, Sudeep Dasari, Siddharth Karamcheti, Soroush Nasiriany, Mohan Kumar Srirama, Lawrence Yunliang Chen, Kirsty Ellis, et al. DROID: A large-scale in-the-wild robot manipulation dataset. *arXiv preprint arXiv:2403.12945*, 2024.

[23] Moo Jin Kim, Karl Pertsch, Siddharth Karamcheti, Ted Xiao, Ashwin Balakrishna, Suraj Nair, Rafael Rafailov, Ethan Foster, Grace Lam, Pannag Sanketi, et al. OpenVLA: An open-source vision-language-action model. *arXiv preprint arXiv:2406.09246*, 2024.

[24] Moo Jin Kim, Chelsea Finn, and Percy Liang. Fine-tuning vision-language-action models: Optimizing speed and success. *arXiv preprint arXiv:2502.19645*, 2025.

[25] Yejin Kim, Wilbert Pumacay, Omar Rayyan, Max Argus, Winson Han, Eli VanderBilt, Jordi Salvador, Abhay Deshpande, Rose Hendrix, Snehal Jauhri, et al. MolmoSpaces: A large-scale open ecosystem for robot navigation and manipulation. *arXiv preprint arXiv:2602.11337*, 2026.

[26] Fuhao Li, Wenxuan Song, Han Zhao, Jingbo Wang, Pengxiang Ding, Donglin Wang, Long Zeng, and Haoang Li. Spatial forcing: Implicit spatial representation alignment for vision-language-action model. In *International Conference on Learning Representations (ICLR)*, 2026.

[27] Haozhan Li, Yuxin Zuo, Jiale Yu, Yuhao Zhang, Zhaohui Yang, Kaiyan Zhang, Xuekai Zhu, Yuchen Zhang, Tianxing Chen, Ganqu Cui, et al. SimpleVLA-RL: Scaling VLA training via reinforcement learning. In *International Conference on Learning Representations (ICLR)*, 2026.

[28] Puhao Li, Yingying Wu, Ziheng Xi, Wanlin Li, Yuzhe Huang, Zhiyuan Zhang, Yinghan Chen, Jianan Wang, Song-Chun Zhu, Tengyu Liu, et al. ControlVLA: Few-shot object-centric adaptation for pre-trained vision-language-action models. *arXiv preprint arXiv:2506.16211*, 2025.

[29] Shuo Liu, Ishneet Sukhvinder Singh, Yiqing Xu, Jiafei Duan, and Ranjay Krishna. VLS: Steering pretrained robot policies via vision-language models. *arXiv preprint arXiv:2602.03973*, 2026.

[30] Mingjie Pan, Siyuan Feng, Qinglin Zhang, Xinchen Li, Jianheng Song, Chendi Qu, Yi Wang, Chuankang Li, Ziyu Xiong, Zhi Chen, et al. SOP: A scalable online post-training system for vision-language-action models. *arXiv preprint arXiv:2601.03044*, 2026.

[31] Karl Pertsch, Kyle Stachowicz, Brian Ichter, Danny Driess, Suraj Nair, Quan Vuong, Oier Mees, Chelsea Finn, and Sergey Levine. FAST: Efficient action tokenization for vision-language-action models. *arXiv preprint arXiv:2501.09747*, 2025.

[32] Stefano Recanatesi, Serena Bradde, Vijay Balasubramanian, Nicholas A. Steinmetz, and Eric Shea-Brown. A scale-dependent measure of system dimensionality. *Patterns*, 3(8):100555, 2022. doi: 10.1016/j.patter.2022.100555.

[33] Mohit Sharma, Claudio Fantacci, Yuxiang Zhou, Skanda Koppula, Nicolas Heess, Jon Scholz, and Yusuf Aytar. Lossless adaptation of pretrained vision models for robotic manipulation. *arXiv preprint arXiv:2304.06600*, 2023.

[34] Octo Model Team, Dibya Ghosh, Homer Walke, Karl Pertsch, Kevin Black, Oier Mees, Sudeep Dasari, Joey Hejna, Tobias Kreiman, Charles Xu, et al. Octo: An open-source generalist robot policy. *arXiv preprint arXiv:2405.12213*, 2024.

[35] Yixiao Wang. Generalization capability for imitation learning. *arXiv preprint arXiv:2504.18538*, 2025.

[36] Mitchell Wortsman, Gabriel Ilharco, Jong Wook Kim, Mike Li, Simon Kornblith, Rebecca Roelofs, Raphael Gontijo Lopes, Hannaneh Hajishirzi, Ali Farhadi, Hongseok Namkoong, et al. Robust fine-tuning of zero-shot models. In *2022 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 7949–7961. IEEE, 2022.

[37] Tian-Yu Xiang, Ao-Qun Jin, Xiao-Hu Zhou, Mei-Jiang Gui, Xiao-Liang Xie, Shi-Qi Liu, Shuang-Yi Wang, Sheng-Bin Duan, Fu-Chao Xie, Wen-Kai Wang, et al. Parallels between vla model post-training and human motor learning: Progress, challenges, and trends. *arXiv preprint arXiv:2506.20966*, 2025.

[38] Youguang Xing, Xu Luo, Junlin Xie, Lianli Gao, Hengtao Shen, and Jingkuan Song. Shortcut learning in generalist robot policies: The role of dataset diversity and fragmentation. *arXiv preprint arXiv:2508.06426*, 2025.

[39] Yajat Yadav, Zhiyuan Zhou, Andrew Wagenmaker, Karl Pertsch, and Sergey Levine. Robust fine-tuning of vision-language-action robot policies via parameter merging. In *International Conference on Learning Representations (ICLR)*, 2026.

[40] Hongzhi Zang, Mingjie Wei, Si Xu, Yongji Wu, Zhen Guo, Yuanqing Wang, Hao Lin, Peihong Wang, Liangzhi Shi, Yuqing Xie, et al. RLinf-VLA: A unified and efficient framework for reinforcement learning of vision-language-action models. *arXiv preprint arXiv:2510.06710*, 2025.

[41] Michał Zawalski, William Chen, Karl Pertsch, Oier Mees, Chelsea Finn, and Sergey Levine. Robotic control via embodied chain-of-thought reasoning. In *Conference on Robot Learning (CoRL)*, 2024.

[42] Lihan Zha, Asher J Hancock, Mingtong Zhang, Tenny Yin, Yixuan Huang, Dhruv Shah, Allen Z Ren, and Anirudha Majumdar. LAP: Language-action pre-training enables zero-shot cross-embodiment transfer. *arXiv preprint arXiv:2602.10556*, 2026.

[43] Wenbo Zhang, Yang Li, Yanyuan Qiao, Siyuan Huang, Jiajun Liu, Feras Dayoub, Xiao Ma, and Lingqiao Liu. Effective tuning strategies for generalist robot manipulation policies. In *2025 IEEE International Conference on Robotics and Automation (ICRA)*, pages 7255–7262. IEEE, 2025.

[44] Xueyang Zhou, Yangming Xu, Guiyao Tie, Yongchao Chen, Guowen Zhang, Duanfeng Chu, Pan Zhou, and Lichao Sun. LIBERO-Pro: Towards robust and fair evaluation of vision-language-action models beyond memorization. *arXiv preprint arXiv:2510.03827*, 2025.

Table 1 | Training configuration for the two stages of RAFT.

	Stage I: reasoning alignment	Stage II: action fine-tuning
Trainable components	LoRA adapters only	Language model, action expert, and projection layers; SigLIP vision tower and Stage-I LoRA frozen
LoRA placement	All 18 language-model layers: q/k/v/o attention projections and gate/up/down MLP projections (126 modules); all 27 vision layers: q/k/v projections (81 modules)	The Stage-I LoRA remains attached and frozen; no adapter parameters are optimized in Stage II
LoRA configuration	Rank $r = 64$, scaling $\alpha = 128$ ($\alpha/r = 2.0$), dropout 0.05	—
Trainable parameters	90.4M (78.4M language-model and 11.9M vision-adapter parameters)	3.09B of 3.50B parameters
Optimizer	AdamW, $\beta = (0.9, 0.999)$, weight decay 10^{-6}, gradient clipping 1.0	AdamW, $\beta = (0.9, 0.95)$, gradient clipping 1.0
Learning rate	Peak 2×10^{-4}; 100 warmup steps; cosine schedule (2,010-step horizon), 1,000 steps	Peak 5×10^{-5}; 1,000 warmup steps; thereafter effectively constant at 5×10^{-5}
Effective batch size	12 samples/GPU × accumulation 4 = 48	16 samples/GPU × 4 GPUs × accumulation 8 = 512
Optimization steps	1,000	5,000
Numerics	bfloat16 with gradient checkpointing	bfloat16 with gradient checkpointing
Hardware	1× NVIDIA L40 (46 GB)	4× NVIDIA L40 (46 GB), one node with distributed data parallel training

Appendix

A Training and implementation details

Table 1 reports the training configuration RAFT uses in simulation and on DROID (SigLIP vision encoder and an 18-layer Gemma-2B decoder). The YAM experiments instead adapt $\pi_{0.5}$ using a similar setup. In RAFT, the pretrained SigLIP weights remain frozen in both stages. Stage I optimizes reasoning LoRA modules in the language model and vision encoder.

A.1 Annotation format

For each observation and task instruction, Gemini 3.1 Pro is asked to return a JSON object in the following format. Coordinates are normalized to the $[0, 1000]$ image frame; boxes are $[x_1, y_1, x_2, y_2]$ and points are (x, y). The format below is for pick-and-place; the open drawer task uses the same structure with `handle_point`, `pre_grasp_point`, `grasp_point`, `pull_start_point`, and `pull_end_point` as keypoints.

```
{
  "reasoning": "<one or two sentences about relevant objects and positions>",
  "objects": [
    {"name": "<label>", "bbox": [x1, y1, x2, y2]}
  ],
  "keypoints": {
```

```
    "grasp_point":     [x, y],
    "target_point":    [x, y],
    "lift_point":      [x, y],
    "pre_grasp_point": [x, y],
    "pre_place_point": [x, y]
  },
  "trace": [[x, y], ...]
}
```

Bimanual (YAM). For the bimanual box-loading task, one arm holds the lid open while the other loads the object, so the keypoints are split by arm rather than by motion stage, and the response adds a per-arm block:

```
{
  "reasoning": "<one or two sentences about relevant objects and positions>",
  "objects": [
    {"name": "<label>", "bbox": [x1, y1, x2, y2]}
  ],
  "keypoints": {
    "left_ee_point":  [x, y],
    "lid_point":      [x, y],
    "right_ee_point": [x, y],
    "object_point":   [x, y],
    "drop_point":     [x, y]
  },
  "trace": [[x, y], ...],
  "arms": {
    "left":  {"motion": "<what the LEFT arm is doing in this frame>",
              "trace": [[x, y], ...]},
    "right": {"motion": "<what the RIGHT arm is doing in this frame>",
              "trace": [[x, y], ...]}
  }
}
```

A.2 Stage-I supervision targets by modality

Table 2 illustrates the target formats for a mug-into-bowl pick-and-place example. Every row uses the same pick-and-place observation and instruction; only the response format changes. VQA and spatial supervision contain language only, while geometric variants use the corresponding coordinates and, where applicable, object or keypoint labels.

A.3 Baseline implementation details

Plain FT. Plain fine-tuning optimizes the SigLIP vision encoder, language model, action expert, and projection layers without Stage-I reasoning alignment.

RETAIN and DeLock. We follow the methods as defined in their original papers [17, 39]. For our evaluations, RETAIN uses the merge coefficient $\alpha = 0.9$ reported by its authors. Both baselines use trainable SigLIP weights during fine-tuning. DeLock regularizes the drift of those weights from their

Table 2 | Modality-specific Stage-I targets for a mug-into-bowl pick-and-place example. Coordinates use the normalized $[0, 1000]$ image frame.

Variant	Target response
VQA	The robot should grasp the glossy beige ceramic mug on the left side of the counter and place it into the white bowl on its right.
Spatial	The beige mug is left of the white bowl. The gripper begins above and to the right of the mug, moves to the mug, lifts it, then moves rightward to the bowl.
Bounding box	`glossy beige ceramic mug: [310, 600, 380, 740]; white bowl: [410, 570, 520, 720]`
Keypoints	`grasp_point: [340, 670]; target_point: [460, 650]; lift_point: [340, 500]; pre_grasp_point: [340, 550]; pre_place_point: [460, 500]`
Trajectory trace	`[[600, 400], [470, 470], [340, 550], [340, 670], [400, 580], [460, 650]]`
Full reasoning	The VQA plan together with the two object boxes, five semantic keypoints, and six-point trajectory trace shown above.

pretrained values with coefficient 10^{-2}. We tune its contrastive-prompt coefficient over the values described below.

Contrastive prompt. We sweep the contrastive-prompt coefficient over $\{0.5, 2, 3\}$ and report the setting with the highest average success across evaluation axes.

A.4 Evaluation protocol

Simulation task definitions, reset procedure, matched initial states, and success criteria follow the MolmoSpaces evaluation protocol [25]. Within each evaluation axis, methods are compared on matched episode IDs. Physical-robot resets preserve the task configuration and start state. DROID used 30 ID trials and 24, 24, and 20 trials for the object, distractor, and background shifts; YAM used 30 trials per OOD axis.

B Qualitative real-robot comparisons

Figures 13–15 compare policies on matched DROID evaluation cases in which RAFT completes the task and the comparison policies fail before completion. Each method is shown at four ordered points in its rollout. Episode names and initial object configurations are matched across methods.

Figures 16–18 show matched YAM evaluations under three OOD shifts. Rows correspond to RAFT, Plain FT, RETAIN, and DeLock; columns follow each rollout in temporal order.

C Additional ablations and controls

Section C.1 varies the format of the Stage-I supervision target; Section C.2 ablates how the Stage-I reasoning adapter is treated during Stage II; Section C.3 compares staged training with co-training and with frozen-vision fine-tuning; and Section C.4 tests whether RAFT composes with RETAIN weight merging.

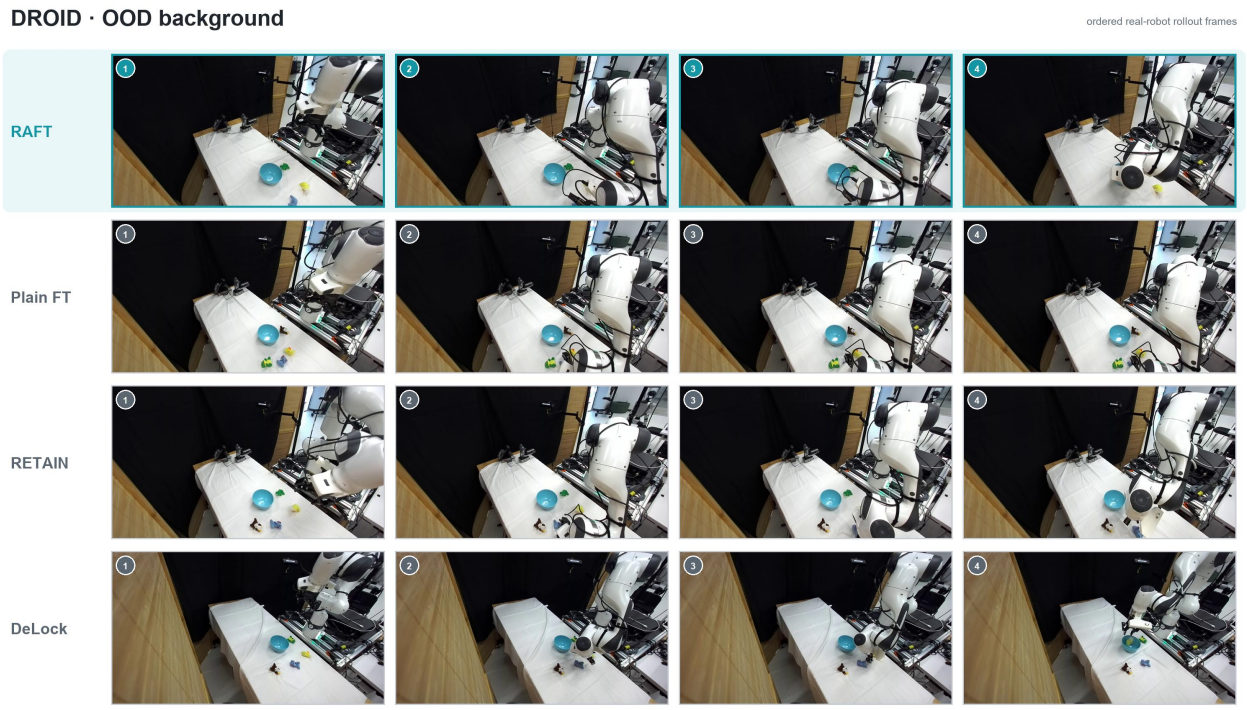

Figure 13 | DROID rollout comparison under an OOD background. Rows show RAFT, Plain FT, RETAIN, and DeLock at four ordered points in matched real-robot episodes. RAFT completes the manipulation, while the comparison policies fail before task completion.

C.1 Effect of the reasoning modality

Setup. We vary the Stage-I supervision target across six formats: full reasoning, visual question answering (VQA), spatial descriptions, bounding boxes, keypoints, and trajectory traces; Table 2 gives an example of each. Each variant changes only the format of the Stage-I target, while Stage-II action training and the evaluation protocol are held fixed. The main-text results use full reasoning for pick and place and bounding-box supervision for the open drawer task, with the modality fixed across evaluation perturbations within each task; the remaining modalities are reported here. Figure 19 summarizes success on both tasks together with a layer-wise participation-ratio analysis.

Results. On the *open drawer* task, trajectory-trace supervision achieves the highest ID success (49.06% vs. 29.25% for Plain FT). Under the visual shift, spatial descriptions, bounding boxes, and trajectory traces reach 29.25%, 28.30%, and 23.58%, respectively. On *pick and place*, full reasoning achieves the highest ID success (95.28%), whereas trajectory traces achieve the highest OOD-object success (45.28%), followed by VQA and full reasoning (43.40% each). Bounding-box and keypoint supervision each reach 34.91%, comparable to Plain FT (35.85%).

C.2 Trainable reasoning adapter in Stage II

Setup. We evaluate the same Stage-I reasoning LoRA under three settings: (i) attached to the base policy without Stage-II action fine-tuning, (ii) updated during Stage II, and (iii) frozen during Stage II, as in RAFT. Table 3 reports pick-and-place success over 106 episodes per split.

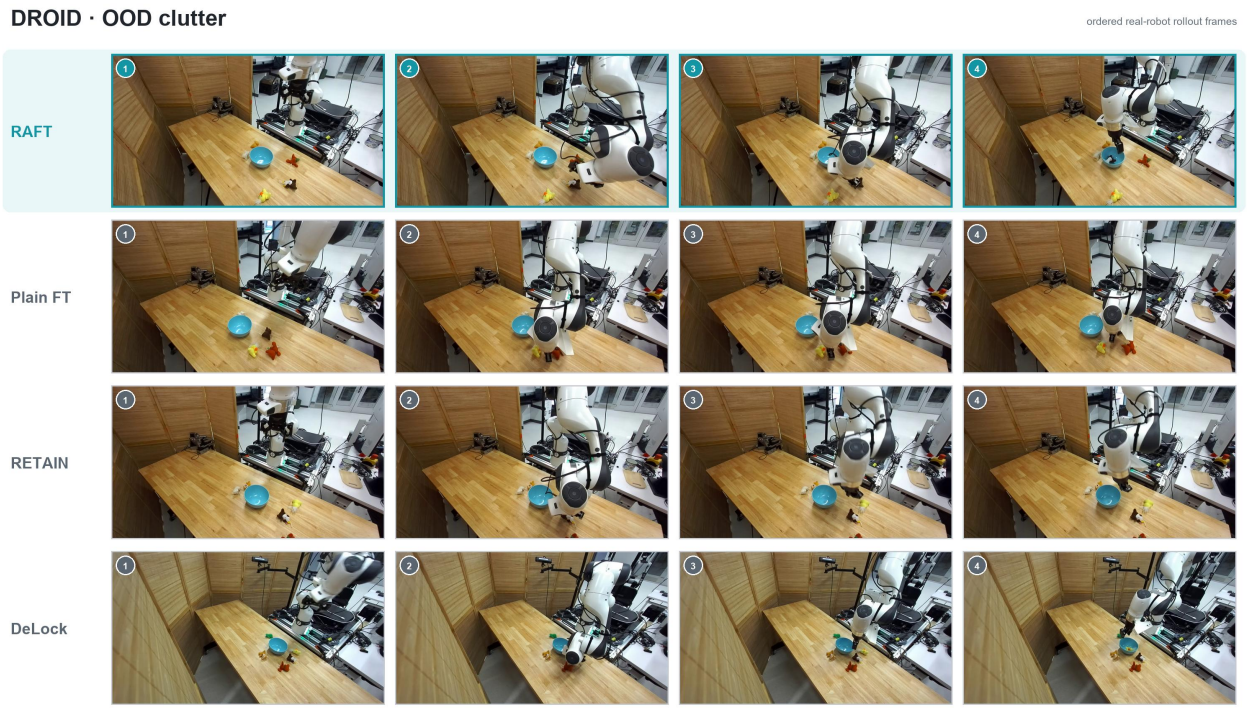

Figure 14 | DROID rollout comparison under OOD clutter. Rows show RAFT, Plain FT, RETAIN, and DeLock at four ordered points in matched real-robot episodes. RAFT completes the manipulation, while the comparison policies fail before task completion.

Table 3 | Stage-II reasoning-adapter ablation on pick and place. Success rate (%) when the Stage-I reasoning LoRA is applied without Stage II, updated during Stage II, or frozen during Stage II (RAFT). Each split contains 106 episodes.

Adapter setting	ID	OOD-light	OOD-texture	OOD-object	OOD-prompt	Mean
LoRA on base (no Stage II)	72.64	60.38	9.43	33.02	66.04	48.30
LoRA trainable in Stage II	90.57	86.79	65.09	34.91	**92.45**	73.96
LoRA frozen in Stage II (RAFT)	**95.28**	**93.40**	**92.45**	**43.40**	91.51	**83.21**

Results. Attaching the LoRA to the base policy without Stage II yields 48.30% mean success. Updating the LoRA during Stage II raises mean success to 73.96%, and freezing it raises it further to 83.21%. The gap between the trainable and frozen settings is largest under the texture shift (65.09% vs. 92.45%).

C.3 Comparison with co-training and frozen-vision fine-tuning

Setup. We compare staged training with co-training, in which reasoning and action supervision are applied within a single training run. Reasoning and action examples are sampled in a 50/50 ratio, and their losses are weighted equally. The primary Plain FT and co-training baselines fine-tune the SigLIP vision encoder. We also evaluate Plain FT with a frozen vision encoder, matching RAFT, which adapts vision only through the Stage-I LoRA and keeps that LoRA frozen in Stage II. Table 4 reports pick-and-place success on the same 106 evaluation episodes per split.

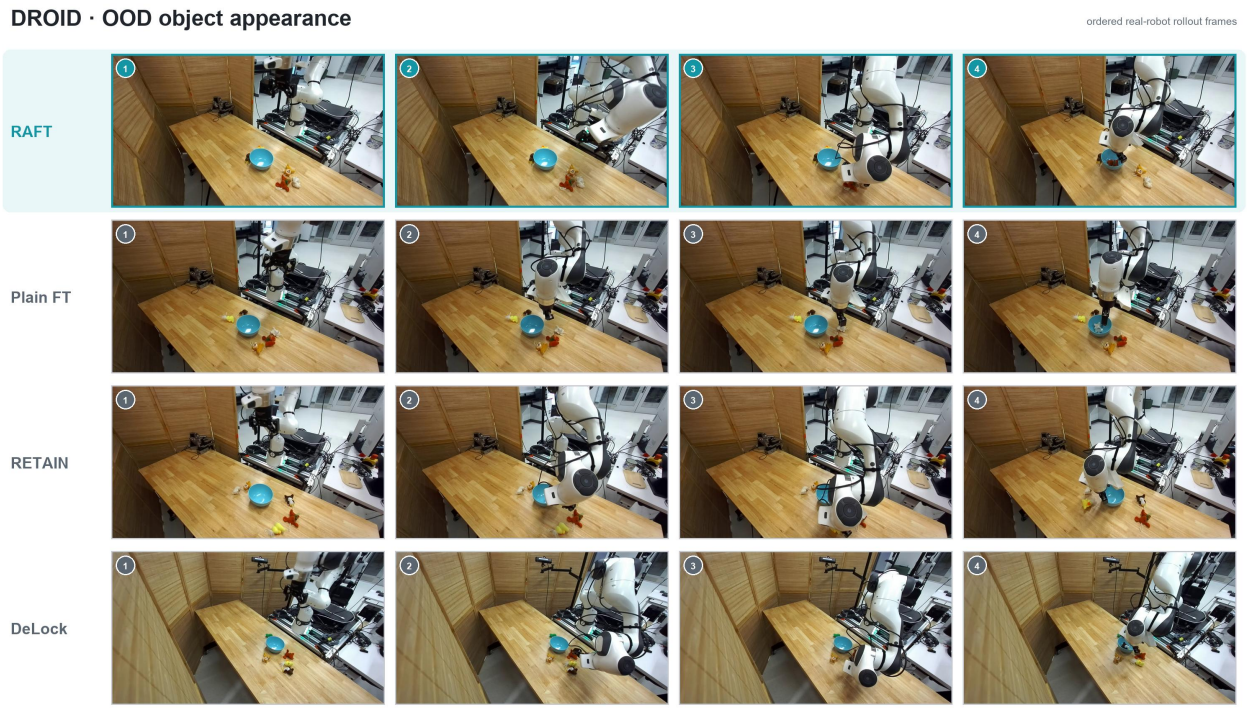

Figure 15 | DROID rollout comparison with an OOD object. Rows show RAFT, Plain FT, RETAIN, and DeLock at four ordered points in matched real-robot episodes. RAFT completes the manipulation, while the comparison policies fail before task completion.

Table 4 | Co-training and frozen-vision controls on simulated pick and place. Success rate (%) on five matched evaluation splits (n = 106 per cell).

Method	ID	OOD-light	OOD-texture	OOD-object	OOD-prompt	Mean
Plain FT	90.57	87.74	**92.45**	35.85	**91.51**	79.62
Plain FT (frozen vision)	91.51	86.79	87.74	40.57	**91.51**	79.62
Co-train	79.25	70.75	72.64	17.92	75.47	63.21
RAFT	**95.28**	**93.40**	**92.45**	**43.40**	**91.51**	**83.21**

Results. Co-training reaches 63.21% mean success, compared with 79.62% for Plain FT and 83.21% for RAFT; its largest deficit is under the object shift (17.92% vs. 43.40% for RAFT). Freezing the vision encoder in Plain FT changes its split-level results but leaves its mean success unchanged at 79.62%. Its OOD-object success is 40.57%, compared with 35.85% for the primary Plain FT and 43.40% for RAFT.

C.4 Compatibility with weight merging

Setup. We test whether RAFT can serve as the adapted checkpoint in RETAIN weight merging. Both merged policies interpolate with the same pretrained base at $\alpha = 0.9$: one uses Plain FT as the adapted endpoint, and the other uses RAFT with bounding-box supervision. We also evaluate the standalone RAFT policy. Table 5 reports results on the Franka open drawer task; every cell contains 106 trials.

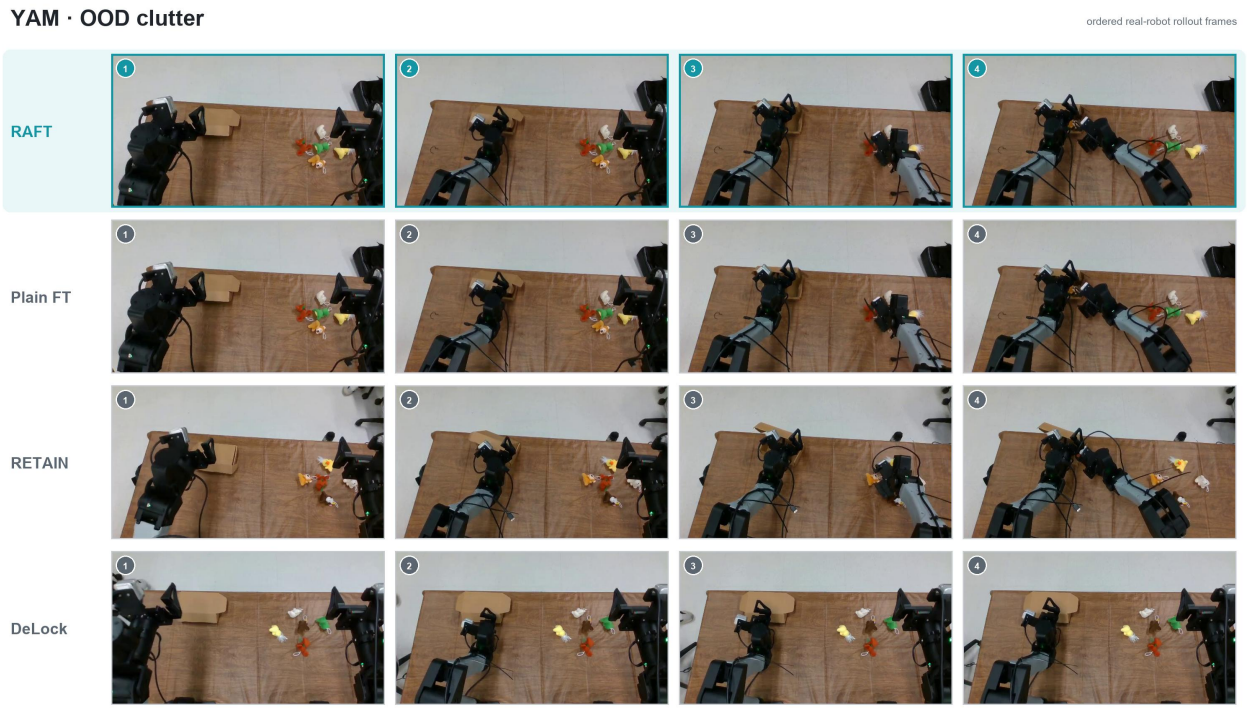

Figure 16 | YAM rollout comparison under OOD clutter. Rows show RAFT, Plain FT, RETAIN, and DeLock on matched task instances; columns show ordered rollout frames. RAFT completes the two-stage task, while the comparison policies do not.

Table 5 | RETAIN applied to Plain FT and RAFT on the Franka open drawer task. Success rate (%) for each axis. RETAIN applied to RAFT outperforms RETAIN applied to Plain FT.

Policy	ID	OOD-light	OOD-texture	Overall
RETAIN on Plain FT ($\alpha = 0.9$)	38.68	34.91	17.92	30.50
RETAIN on RAFT (bbox, $\alpha = 0.9$)	41.51	37.74	**29.25**	36.16
RAFT (bbox) alone	**44.34**	**41.51**	28.30	**38.05**

Results. Using RAFT instead of Plain FT as the adapted endpoint increases pooled success from 30.50% to 36.16%. Standalone RAFT reaches 38.05%, the highest overall, with the highest ID and lighting success (44.34% and 41.51%).

D Evaluation environments

Figure 20 illustrates the in-distribution (ID) and out-of-distribution (OOD) conditions used in our simulated and physical-robot experiments; the evaluation protocol is described in Appendix A.4. Simulation task construction, resets, initial-state matching, and success criteria follow MolmoSpaces [25]. Within each perturbation, all methods are evaluated on identical episodes. On the physical robots, resets preserve the task configuration and initial state.

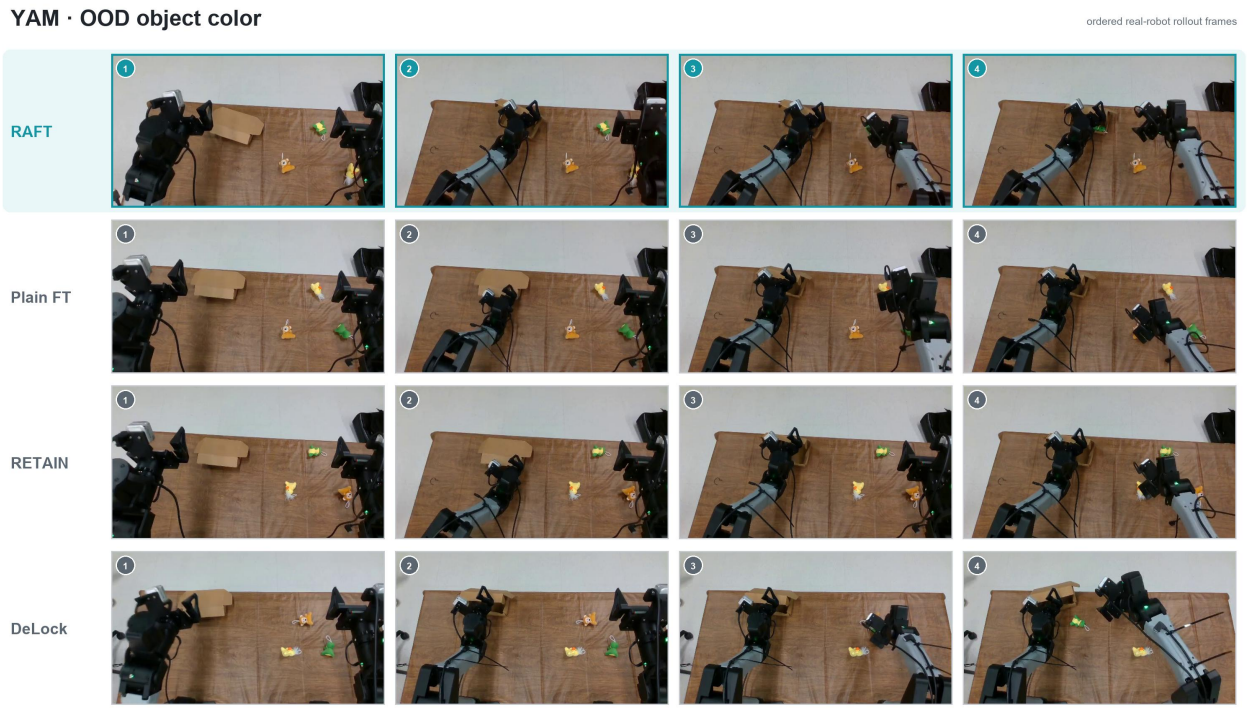

Figure 17 | YAM rollout comparison under an OOD color shift. Rows show RAFT, Plain FT, RETAIN, and DeLock on matched task instances; columns show ordered rollout frames. RAFT completes the two-stage task, while the comparison policies do not.

E Representation and attention analysis protocol

Model and captured frames. The participation-ratio (PR) analysis in Figure 19 uses the first 400 records of the held-out iTHOR FloorPlan13 validation split, spanning five episodes. The linear-probe results in Figure 9 use a capture of 414 frames: nine evenly spaced frames from each of 46 policy-training episodes.

Participation ratio and correlations. For each frame and decoder layer ℓ, we mean-pool `hidden_states[`ℓ`]` over non-padding positions to obtain a 400×2048 matrix X_ℓ. After centering the columns, we compute the covariance eigenvalues λ_i from the squared singular values and define $\mathrm{PR}_\ell = (\sum_i \lambda_i)^2 / \sum_i \lambda_i^2$. PR thus measures the effective dimensionality of *across-frame* variance. Each evaluation axis pairs one downstream benchmark with one PR probe domain. Figure 19 pools the three task-matched axes. For each layer and axis, we compute the Spearman correlation across the same six Stage-I arms and assess significance with a two-sided exact permutation test over all $6! = 720$ assignments, applying Benjamini–Hochberg FDR correction across the 19 layers.

Linear probes. Episode-time labels are defined as $\min(3, \lfloor 4\,\mathrm{frame_idx}/\mathrm{total_frames} \rfloor)$, giving a chance balanced accuracy of 1/4. The labels *reach*, *grasp*, *transport*, and *place* denote the four positional (timestep) quartiles of each episode. Within each leave-one-episode-out fold, we fit `StandardScaler`, `PCA(64)`, and `LogisticRegression(C=0.5)`. Layer-band uncertainty and differences from Plain FT are estimated with a paired bootstrap over episodes ($B = 1000$).

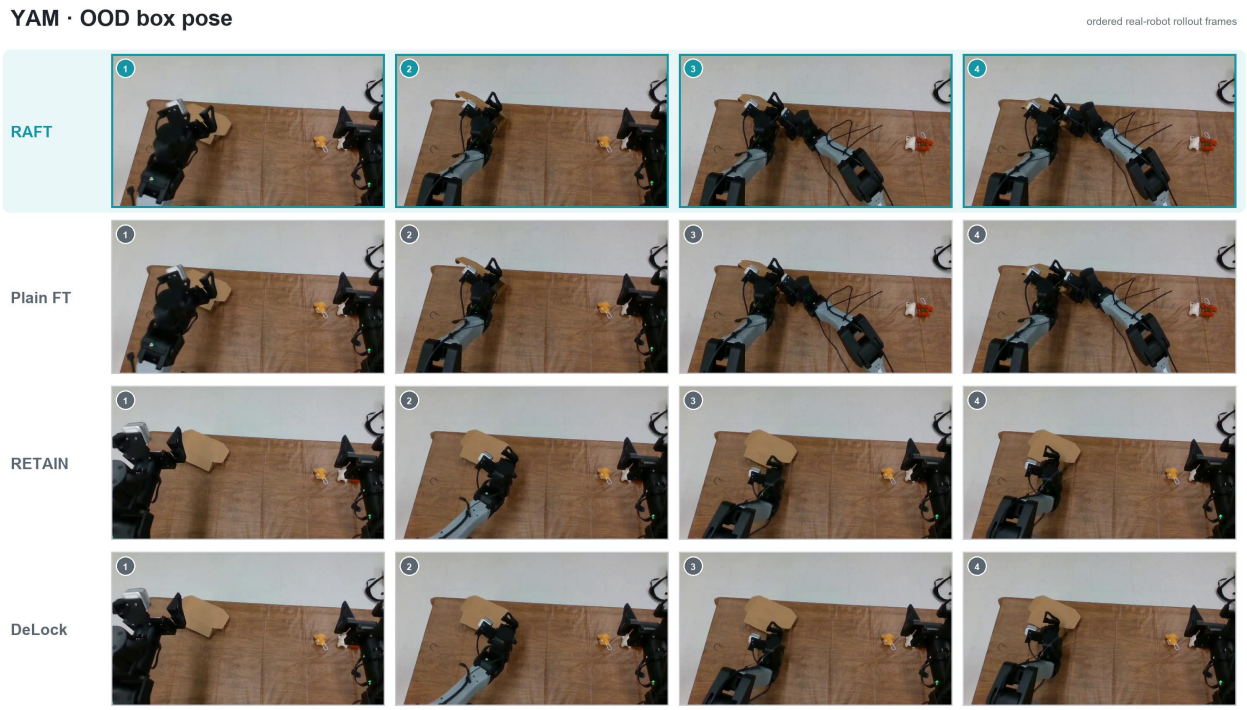

Figure 18 | YAM rollout comparison under an OOD initial pose. Rows show RAFT, Plain FT, RETAIN, and DeLock on matched task instances; columns show ordered rollout frames. RAFT completes the two-stage task, while the comparison policies do not.

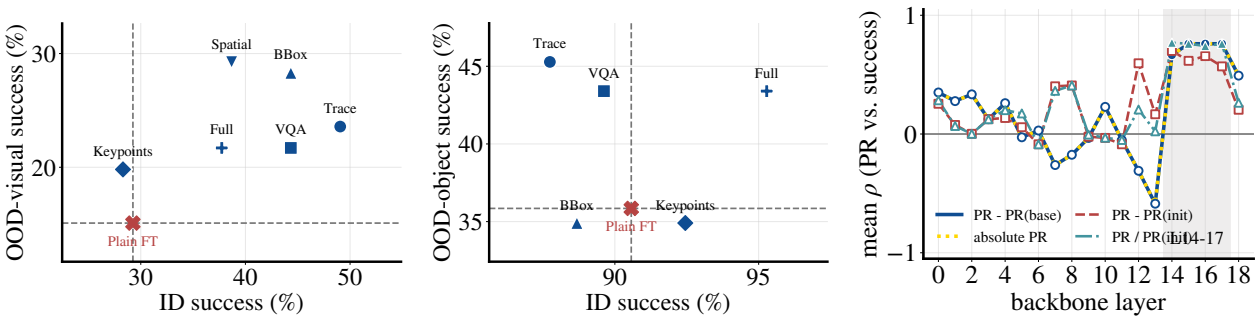

Figure 19 | Effect of reasoning modality on success and representation dimensionality. *Left:* ID and OOD-visual success on the open drawer task for six reasoning modalities. *Center:* ID and OOD-object success on pick and place for five modalities. Red crosses denote no-reasoning controls, and dashed lines mark Plain FT. Each policy is evaluated on 106 trials per axis. *Right:* Layer-wise mean Spearman correlation between participation ratio and success, pooled over three task-matched evaluation axes and six reasoning variants, for four participation-ratio measures. Filled markers denote significance after FDR correction; shading marks layers 14–17.

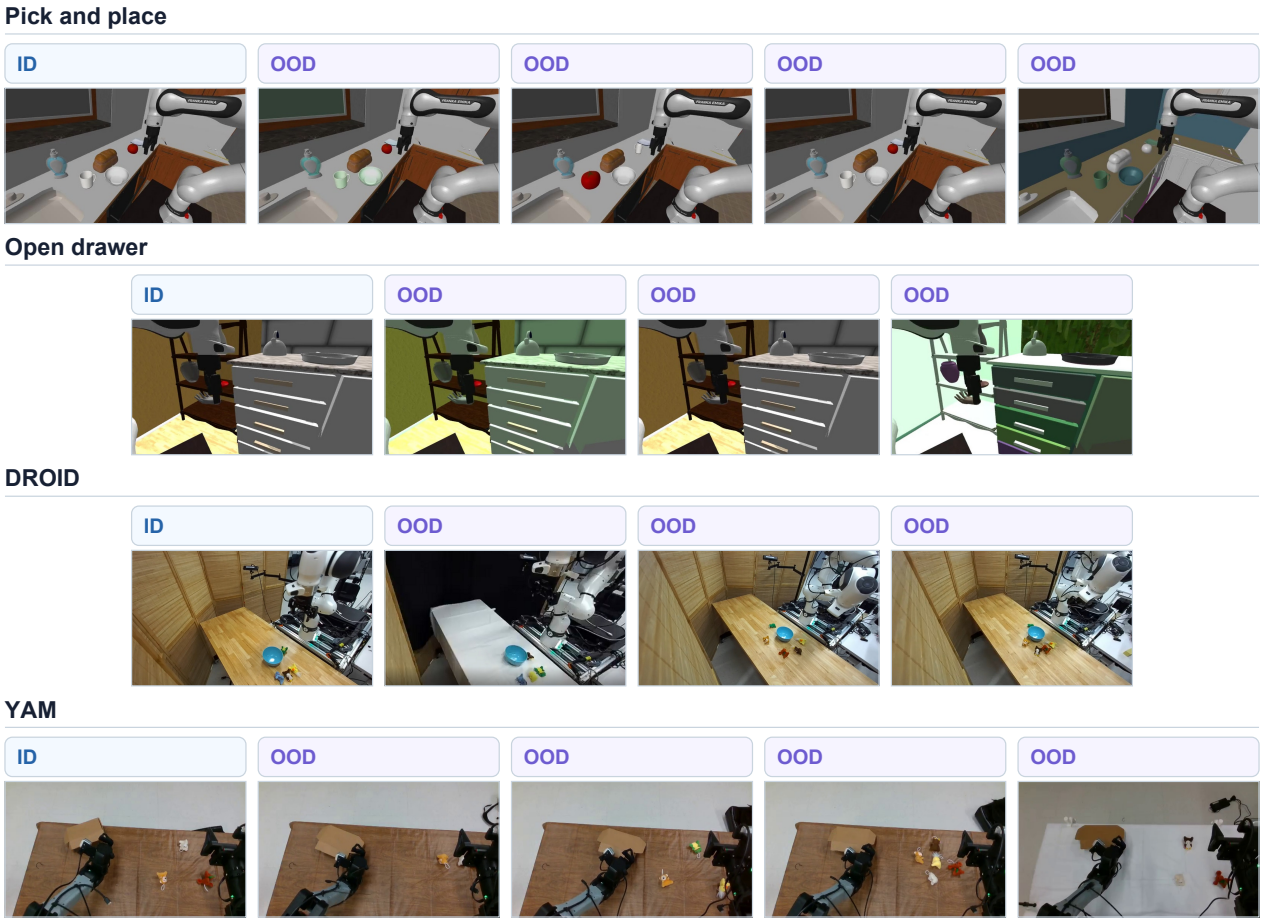

Figure 20 | Evaluation setup across the four robot tasks. Each row begins with an in-distribution observation, followed by its OOD variants. Pick and place varies lighting, target object, instruction, and scene texture; the open drawer task varies lighting, instruction, and scene texture; DROID varies background, clutter, and object appearance; and YAM varies box pose, object color, clutter, and background. Instruction shifts reuse the ID image because only the task prompt changes.

Visual attention and region masks. At each layer, we extract the attention from text queries to image keys, $A[:, 256:, :256]$, average it uniformly over heads and text queries, and L_1-normalize it over the 256 image patches; no individual head or last-token readout is selected. Band-level maps average the normalized per-layer maps. Because this normalization removes the total attention mass on the image, we report that quantity separately as `img_mass`. For visualization, patch values are upsampled by nearest-neighbor tiling, and all arms share a color scale capped at the 99th percentile of interior-patch attention; the border ring is excluded from this scale because it exhibits attention-sink/register behavior. Bounding boxes are mapped to patches using floor/ceiling boundaries and may include up to one extra patch on each edge; each keypoint or trace point is mapped to the patch that contains it, and the background is the complement of all annotated patches. A frame enters a given region analysis only if it contains at least two bounding-box patches, one keypoint, or one trace point, as applicable. For a patch mask M and normalized attention p_j, we define enrichment as $(256/|M|)\sum_{j\in M} p_j$, so that a value of 1 corresponds to uniform attention. Bounding-box, keypoint, and trace masks may overlap and therefore do not partition the image; only the background mask is defined by complementation.

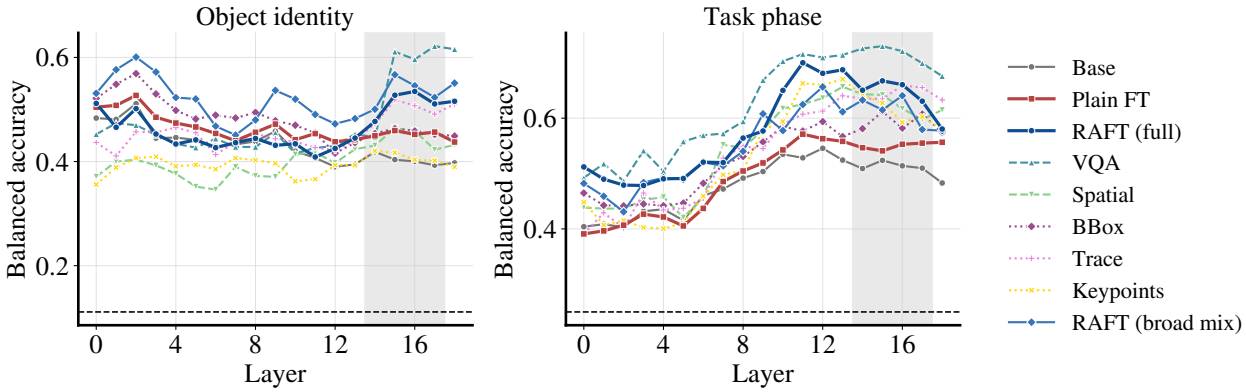

Figure 21 | Late-layer probes for all nine model arms. Leave-one-episode-out balanced accuracy for object identity and episode-time quartile across backbone layers. The shaded band marks layers 14–17. This expands the four-arm comparison in the main text to the full set of reasoning-modality variants.

F Extended late-layer probe analysis

Figure 21 expands the main-text probe comparison to all nine model arms. VQA and trace show the largest late-layer recovery in object-identity decoding and the strongest novel-object results; spatial, bounding-box, and keypoint variants show no comparable recovery and are the weakest on that shift.