arXiv:2611.14380v1 [cs.CV] 10 Nov 2026

What Factors Matter for Collection-Level Counting with Human–AI Collaboration?

Seoyun Jeong[1], **Logan Lawrence**[1], **Owen L. Martin**[2], **Benjamin Koger**[2], **Daniel Sheldon**[1], **Subhransu Maji**[1]

[1]University of Massachusetts Amherst
[2]University of Wyoming
seoyunjeong@umass.edu, lclawrence@umass.edu, omartin9@uwyo.edu, bkoger@uwyo.edu, sheldon@cs.umass.edu, smaji@umass.edu

Abstract

Many scientific and policy decisions require estimating the total number of target objects across a large image collection. Exhaustive human labeling is often infeasible, while raw model predictions can be biased, especially under domain shift. Recent human–AI estimation methods combine imperfect model predictions with limited human verification to obtain reliable collection-level estimates. However, existing work provides little guidance on how practitioners should choose among models, estimators, and adaptation strategies for a new deployment. We study these choices within a common human-in-the-loop counting framework across 12 tasks spanning biomedical microscopy, ground-level and aerial wildlife surveys, and overhead imagery of urban structures. Our results show that modern multimodal models provide a strong, training-free starting point for many tasks, while fine-tuned detectors become more advantageous at larger labeling budgets and across repeated deployments. Estimation efficiency improves further when model predictions are used to prioritize which units humans label, rather than sampling uniformly. We also find that each model's ranking quality can be estimated with a small pilot set, yielding practical guidance for selecting strategies on new counting tasks. In a field deployment on aerial pronghorn surveys, a task-specific adaptation of these methods meets the survey's $\pm 5\%$ tolerance after reviewing a median of only 3.2% of collected images.

1 Introduction

Many scientific and policy questions involve counting in image collections: How many infected cells are present in a set of blood smears? How many birds migrate through a region in a given year? How many buildings were damaged in a disaster-affected region? Such totals can inform public-health surveillance (Rajaraman et al. 2018), conservation policy (Norouzzadeh et al. 2018; Eikelboom et al. 2019), and disaster response (Gupta et al. 2019), aiding allocation of limited resources. However, in each case, the imagery needed to answer these questions is abundant, while exhaustive human review does not scale to the collection sizes now routinely produced by their imaging pipelines. Similar problems arise in many other social impact domains where practitioners often have limited resources for data labeling and model development but require high precision or statistical guarantees that are not easily achieved by pure AI methods.

Figure 1: Samples from the 12 counting tasks. Each in-image label identifies the task (dataset/class), tile index, and exact ground-truth count. Red points mark each object.

A growing body of work combines imperfect model predictions with a small amount of human-verified labels to obtain statistically valid estimates of population totals (Meng et al. 2022; Angelopoulos et al. 2023; Perez, Maji, and Sheldon 2024; Hamilton et al. 2026b). The high-level idea is to use the model to make predictions on the entire collection, then sample images for human labeling to correct the model's total within a statistical estimation framework. These methods can provide unbiased estimates and confidence intervals regardless of the predictor's accuracy. However, their *efficiency*—how much human labeling is needed to achieve a desired level of precision—depends strongly on the predictor model, estimator, data distribution, and labeling protocol.

Prior work offers limited practical guidance on selecting human–AI counting strategies under different deployment constraints. Practitioners may begin with no training data, a detector trained on a related domain, or sufficient resources to train and adapt a custom model. These settings create trade-offs among model quality, labeling efficiency, computational cost, and engineering effort. We therefore ask: Given a target image collection and a set of resource constraints, which model and estimator should one use, and when is adaptation worthwhile? Alongside a new benchmark, we demonstrate

the practical value of this model-assisted counting framework on real-world aerial pronghorn surveys, where counting from imagery rather than physically from the aircraft removes the most dangerous task wildlife professionals perform (Sasse 2003). Our contributions are:

- A benchmark for collection-level counting with 12 tasks across biomedical microscopy, ground-level and aerial wildlife imagery, and overhead urban imagery (Fig. 1).
- A controlled comparison of models, estimators, and adaptation strategies under a common human-in-the-loop estimation framework, together with a procedure for pilot-based model selection.
- A case study on aerial pronghorn surveys that puts these findings into practice, including a lightweight adaptation scheme for repeated deployments.

In preview, we find that modern off-the-shelf open-source multimodal large language models are a remarkably strong starting point for model-assisted counting. On nearly every task in our benchmark, they reach counting error below 10% with under 10% of the collection labeled, cutting error by roughly $4\times$ relative to a model-free estimate that uses only the human-labeled units, at the same budget. Training a custom detector can lower the error further, but at the cost of larger labeling budgets and model training time. On the pronghorn survey, $\pm 5\%$ uncertainty tolerance is achieved after reviewing a median of only 3.2% of the collected images.

2 Related Work

Counting objects in images has been studied for well over a decade. Detection-based methods localize and then add up each instance using general object detectors such as Faster R-CNN (Ren et al. 2015), while density-based methods instead regress a per-pixel density whose integral gives the count, sidestepping explicit detection in crowded scenes (Lempitsky and Zisserman 2010; Zhang et al. 2016). More recently, class-agnostic models extend to novel categories from a few exemplars (Ranjan et al. 2021) or from text queries, as in open-vocabulary detectors (OVDs) and multimodal large language models (MLLMs), with a growing set of benchmarks measuring their reliability (Paiss et al. 2023; Li et al. 2025). Across these paradigms, evaluation targets accuracy on individual images, rather than a total across an entire collection. We instead benchmark models by how well they support estimating collection-level totals, a criterion that does not always track per-image counting accuracy, as our experiments show.

A complementary line of work pairs imperfect model predictions with a small number of human-verified labels to obtain statistically valid estimates of a collection-level total. Its foundations lie in survey sampling, where auxiliary variables correlated with the target yield ratio, regression, and other model-assisted estimators (Cochran 1977; Särndal, Swensson, and Wretman 2003). Recent methods replace hand-chosen covariates with learned predictors: ISCount (Meng et al. 2022) and DISCount (Perez, Maji, and Sheldon 2024) sample images by importance weights derived from model scores, and active measurement (Hamilton et al. 2026b) interleaves labeling, estimation, and model adaptation in a sequential loop. Prediction-powered inference (Angelopoulos et al. 2023) extends the same principle to general estimands, correcting black-box predictions with a labeled subset to form valid confidence intervals, and the related active-testing literature draws informative samples to evaluate models label-efficiently (Kossen et al. 2022; Farquhar, Gal, and Rainforth 2021). Our work builds on advances in both counting models and statistical estimators. We evaluate both fine-tuned and frozen models, and consider practical factors such as inference costs and required software expertise to adapt models to target distributions.

3 Estimating Collection-Level Counts

3.1 Problem Setup

We study estimating a single-class *collection-level total*, i.e., the number of target objects across an image collection. The collection is represented as a finite pool Ω of $N = |\Omega|$ units. Each unit $s \in \Omega$ is a cropped image tile or, for images too small to tile, the whole image (§4.1). Let $f(s) \in \mathbb{Z}_{\geq 0}$ denote the ground-truth count in unit s and $F(S) = \sum_{s \in S} f(s)$ be the total of any subset $S \subseteq \Omega$. The target is $F_{\text{true}} \equiv F(\Omega) = \sum_{s \in \Omega} f(s)$. We formulate this problem within Hamilton et al. (2026b)'s active measurement framework, in which true per-unit counts are acquired sequentially and combined with model predictions to estimate a collection-level total. It generalizes DISCount (Perez, Maji, and Sheldon 2024) to allow model adaptation and sampling without replacement. These features translate well to deployment settings where practitioners accumulate labels and never label the same unit twice.

Assume that a model provides a predicted count $g(s) \geq 0$ for each unit. Before the main labeling loop, fine-tuned detectors are conditioned with a labeled warmup set $\mathcal{W}$, where at least half contain one or more instances of the target object. For frozen models, $\mathcal{W} = \emptyset$. At time steps $t = 1, 2, \ldots, T$, let $\mathcal{D}_t = \mathcal{W} \cup \{s_1, \ldots, s_{t-1}\}$ denote the units revealed before the step. A sampler draws an unlabeled unit $s_t \sim q_t$ from the remaining pool $\mathcal{R}_t = \Omega \setminus \mathcal{D}_t$, where q_t is the current proposal distribution over the remaining pool. Then $f(s_t)$ is revealed, corresponding to a human annotation. Next, an estimator combines the true counts on $\mathcal{D}_t^+ = \mathcal{D}_t \cup \{s_t\}$ with the model predictions g_t available at step t, where $g_t(s)$ is the predicted count on unit s, to produce an estimate $\hat{F}$ of the total. For frozen models, $g_t = g_0$ for all t. Finally, $\mathcal{D}_t^+$ is optionally used to fine-tune the model. To summarize, a counting strategy consists of (i) a prediction model that produces g_t, (ii) a proposal distribution (i.e., sampler) q_t, (iii) an estimator for the total, and (iv) a model-update policy.

3.2 Models

The framework requires only a model that produces a per-unit count; any such model can be substituted. We consider two groups of models a practitioner might choose for a new task: detectors fine-tuned on the target data, and larger frozen models used off the shelf.

Fine-tuned detectors We consider two Faster R-CNN detectors (Ren et al. 2015), one with a larger ResNet-50 (He

et al. 2016) backbone and one with a smaller MobileNetV3-Large (Howard et al. 2019) backbone. Both are initialized from COCO-pretrained weights and adapted to a single foreground class by replacing the box-prediction head with a binary object-versus-background head. They are initially fine-tuned on the warmup set of $w = 50$ units, and continually retrained every 25 steps during the sequential labeling loop. Training details are provided in Appendix F.

Frozen OVDs and MLLMs Given an input image and a target object name, OVDs return localized detections. We use the number of detections above a 0.5 confidence threshold as the per-unit count. We evaluate three OVDs: Florence-2 (Xiao et al. 2024), Grounding DINO (Liu et al. 2024), and OWLv2 (Minderer, Gritsenko, and Houlsby 2023).

MLLMs take an image and a natural-language prompt and respond in free-form text. For all MLLMs, we use the prompt *"How many {object_name} are in this image?"* and extract the count using a constrained LLM-based answer extractor (Lawrence et al. 2026). Predictions without a clear numeric answer are assigned a count of zero. To compare locally runnable open-source models with proprietary models, we evaluate five open-source models (Qwen3-VL-8B (Bai et al. 2025), InternVL3.5-8B (Wang et al. 2025), Molmo-7B-D-0924 (Deitke et al. 2025), Molmo2-8B (Clark et al. 2026), Aya Vision-8B (Dash et al. 2025)) and three commercial models (GPT-5.2 (OpenAI 2025), Gemini 3.1 Pro (Google DeepMind 2026), Qwen3.7-Plus (Qwen 2026)). Model-specific inference and cost details are in Appendix I.

3.3 Estimators and Sampling Distribution

We compare four ways of using the model predictions g_t: trust them on the unlabeled pool, ignore them, correct them with uniformly sampled units, or use them to direct label acquisition. All sampling is without replacement. Except for the model-only baseline, the estimators are unbiased (see Appendix A). Our estimators use two different strategies: *batch* estimators use the entire sampled set $\mathcal{D}_t^+$ in each step to form an estimate $\hat{F}_{1:t}$; *incremental* estimators form separate estimates $\hat{F}_t$ in each step and later combine per-step estimates into an overall estimate. Below, we derive expressions for the total; analogous expressions for the variance can be used to construct confidence intervals (Owen 2013; Appendix A).

Model-only baseline *Can we trust the model alone?* For a fair sequential comparison, the model-only baseline replaces predictions on revealed units with their exact counts and trusts the active model on the remainder. Writing $G_t(S) = \sum_{s\in S} g_t(s)$, the batch estimator is:

$$\hat{F}_{1:t}^{\text{model}} = F(\mathcal{D}_t^+) + G_t(\Omega \setminus \mathcal{D}_t^+). \tag{1}$$

Kossen et al. (2022) adopt the same approach, which they call model-based *interpolation*, for active testing.

Model-free Monte Carlo baseline *What can we get from labels alone?* The model-free Monte Carlo baseline is a batch estimator that follows the classical expansion estimator under simple random sampling (Cochran 1977):

$$\hat{F}_{1:t}^{\text{MC}} = F(\mathcal{W}) + \frac{N-w}{t}\sum_{i=1}^{t} f(s_i). \tag{2}$$

The second term scales the mean of the sampled units to estimate the total of the $N - w$ non-warmup units.

Model-assisted uniform sampling estimators *Can uniform samples correct the model's total?* Let $M_t = |\mathcal{R}_t|$. We consider the incremental estimator family

$$\hat{F}_t^{(c_t)} = F(\mathcal{D}_t) + c_t G_t(\mathcal{R}_t) + M_t[f(s_t) - c_t g_t(s_t)] \tag{3}$$

parameterized by coefficient c_t, where both g_t and c_t are fixed before s_t is drawn uniformly from $\mathcal{R}_t$. So they may depend on everything revealed so far, but not on s_t. We instantiate three members of this family, the *incremental* difference, ratio, and regression estimators, which follow classical estimators by the same names (Särndal, Swensson, and Wretman 2003) but are modified to be exactly unbiased under coefficients c_t and model sequences g_t that adapt to revealed samples. The incremental difference estimator uses $c_t = 1$, the ratio estimator uses $c_t = \hat{\gamma}_t = F(\mathcal{D}_t)/G_t(\mathcal{D}_t)$, and the regression estimator uses

$$c_t = \hat{\beta}_t = \frac{\sum_{s\in\mathcal{D}_t}(g_t(s) - \bar{g}_t)\left(f(s) - \bar{f}_t\right)}{\sum_{s\in\mathcal{D}_t}(g_t(s) - \bar{g}_t)^2},$$

where $\bar{g}_t$ and $\bar{f}_t$ are means over $\mathcal{D}_t$. To the best of our knowledge, these estimators are novel. The classical ratio and regression estimators fit their coefficients on the full sample, which makes them biased; computing $\hat{\gamma}_t$ and $\hat{\beta}_t$ from the previously revealed $\mathcal{D}_t$ and holding them fixed for the fresh draw removes that bias. The final term is then an unbiased estimate of the residual total on $\mathcal{R}_t$ at every step and under any sampling history, so all three estimators are unbiased for F_{true} regardless of the accuracy of g_t (Appendix A). Degenerate or under-determined coefficient fits revert to $c_t = 1$.

Model-assisted importance sampling estimators *Can the model choose which unit to label next?* Importance sampling (IS) uses model predictions to prioritize units expected to contribute more to the collection total. DISCount implements this idea with a fixed proposal (DIS; Perez, Maji, and Sheldon 2024), while active measurement (Hamilton et al. 2026b) extends to sampling without replacement and an adaptive proposal (known as *adaptive importance sampling* or AIS; Owen 2013) via the incremental estimator

$$\hat{F}_t^{\text{IS}} = F(\mathcal{D}_t) + f(s_t)/q_t(s_t). \tag{4}$$

Because q_t is likewise fixed before the draw, the second term is an unbiased estimate of the remaining total whenever $q_t(s) > 0$ for every $s \in \mathcal{R}_t$. The estimator has minimum variance (of zero) when $q_t(s) \propto f(s)$, which motivates sampling in proportion to the model's predicted counts. We evaluate four proposal strategies that share this estimator: AIS, DIS, spotAIS, and spotDIS. Define

$$q_t(s) = \frac{\max\{g_t(s), \epsilon\}}{\sum_{u\in\mathcal{R}_t}\max\{g_t(u), \epsilon\}} \tag{5}$$

where $\epsilon > 0$ is used to ensure $q_t(s) > 0$. AIS is the case where g_t, and hence q_t, is time-varying (i.e., the model is retrained), and DIS is the case where $g_t = g_0$ for all t (i.e., the model is not retrained). AIS and DIS use $\epsilon = 1$ following

Estimator	Sampler	Model
Monte Carlo	uniform	—
{model-only, diff, ratio, reg}	uniform	{fine-tuned, frozen}
AIS, spotAIS	q_t	fine-tuned
DIS, spotDIS	q_t	frozen

Table 1: Estimator–sampler–model configurations evaluated. Uniform-sampling estimators are run with both fine-tuned and frozen models, while the IS variants are model-specific, with AIS paired to fine-tuned detectors and DIS to frozen models. Here diff, ratio, and reg denote the incremental difference, ratio, and regression estimators, defined in §3.3.

prior work. We also introduce two new "spot-checking" variants, spotAIS and spotDIS, that ensure positivity through a different mechanism. They allocate a fraction λ of the sampling budget to "spot check" units predicted to be empty and spend the rest in proportion to the model's counts:

$$q_t = \lambda\, q_t^{\text{spot}} + (1-\lambda)\, q_t^{\text{count}} \tag{6}$$

where q_t^{spot} is uniform over predicted-empty units, and q_t^{count} is Eq. (5) with $\epsilon = 0$ restricted to the predicted-populated units; we use $\lambda = 0.05$. The two components have disjoint support, so every unit draws its probability from exactly one term. A unit is predicted empty when $g_t(s) = 0$ or a lightweight binary logistic regression classifier h_t flags it as empty. Trained on frozen DINOv3 (Siméoni et al. 2025) embeddings of the units in $\mathcal{D}_t$, h_t refits in milliseconds and adds a representation-based signal beyond the detector's count. See Appendix B for more details.

Combining per-step estimates For incremental estimators, each main-loop draw yields an unbiased per-step estimate of F_{true}, whether from $\hat{F}_t^{(c_t)}$ or from $\hat{F}_t^{\text{IS}}$. They are combined to form an overall estimate $\hat{F}_{1:t} = \sum_{i=1}^{t} \bar{\alpha}_{i,t} \hat{F}_i$ using weights $\bar{\alpha}_{i,t} > 0$ that sum to one for each t. Because these weights are fixed in advance and never depend on the observed labels, $\hat{F}_{1:t}$ is unbiased as well (Appendix A). We use the "COMB" weighting scheme (Hamilton et al. 2026b) when models are fine-tuned, which gives more weight to later estimates based on two different factors: an improving model and a smaller remaining pool. We alternatively use the "LURE" weighting scheme when models are not updated (Farquhar, Gal, and Rainforth 2021). Details are in Appendix C. The model-only and Monte Carlo baselines are batch estimators and do not require combination.

Table 1 summarizes the counting strategies we evaluate, represented by each estimator–sampler–model triple.

4 Benchmark Setup

4.1 Datasets and Task Construction

We build the benchmark from four publicly available, exhaustively annotated detection datasets spanning distinct visual domains. *Malaria Cell Images (BBBC041v1)* (Ljosa, Sokolnicki, and Carpenter 2012) contains microscopy images annotated with seven cell categories: six states of malaria infection and a *difficult* category for cells that cannot be confidently typed. *Birds* (Hamilton et al. 2026b) consists of image tiles from two high-resolution ground-level photographs of Tree Swallow flocks. *Kenyan Wildlife Aerial Survey* (Eikelboom et al. 2019) comprises aerial survey photographs annotated for elephants, giraffes, and zebras. *DOTA v1.5* (Xia et al. 2018) contains aerial and satellite imagery with annotations for sixteen categories of urban objects, such as planes and storage tanks. Each dataset-class pair defines a counting task. For each task, we collect images containing the target object from the dataset and convert them into a set of estimation units. Fig. 1 shows examples of each task.

Tier (%∅)	Dataset	Class	%∅	μ	N
full ($< 5\%$)	Malaria	trophozoite	0.0	2.42	655
	Malaria	difficult	0.0	1.29	347
	Malaria	red blood cell	2.1	49.19	1,688
partial (5–60%)	Birds	bird	24.4	14.36	1,302
	DOTA	tennis court	37.5	4.63	689
	DOTA	harbor	52.1	3.90	2,084
sparse (60–90%)	DOTA	ship	78.7	2.10	5,061
	DOTA	plane	79.4	0.67	5,043
	DOTA	small vehicle	80.8	1.31	5,032
rare ($\geq 90\%$)	Wildlife	giraffe	91.7	0.16	5,005
	DOTA	helicopter	92.8	0.47	1,523
	DOTA	storage tank	98.1	0.11	5,176

Table 2: The 12 benchmark tasks. %∅ is the empty fraction, μ the average object count per unit, and N the total number of units. Full per-task statistics are in Appendix E.

To standardize apparent object scale, we estimate a task-specific unit size from 30 annotated objects sampled from roughly ten whole images, choosing the resolution so that the median object appears at about 80 pixels after resizing the tile to the Faster R-CNN's standard resize target, 800 pixels. Images whose longer side is at most 1.5× the tile size are kept whole to avoid thin edge tiles; larger images are split into non-overlapping square tiles, cropped at the right and bottom edges. These 30 samples are used only for tiling configuration and do not count toward the labeling budget or warmup. Objects are assigned to the tile containing the center of their bounding box, preventing double counting and ensuring that counts sum to F_{true} (details in Appendix D).

We characterize each task by its *empty fraction*, defined as the proportion of units containing no target objects, to study how object sparsity in the collection relates to different counting strategies. Across the four datasets, 27 classes are grouped into four occupancy tiers based on this fraction: *full* ($< 5\%$), *partial* (5–60%), *sparse* (60–90%), and *rare* ($\geq 90\%$). From each tier, we select the three classes with the largest F_{true} while preserving diversity in object type, yielding twelve tasks (Table 2).

4.2 Evaluation Protocol and Metrics

We evaluate every task and counting strategy combination by the absolute fractional error $\text{AFE}_t = \left|\hat{F}_{1:t} - F_{\text{true}}\right| / F_{\text{true}}$ averaged over 1,000 trials and traced from $t = 1$ to $t = N/2$. Reporting the whole curve exposes how fast each strategy converges and where strategies cross over. Frozen-model tri-

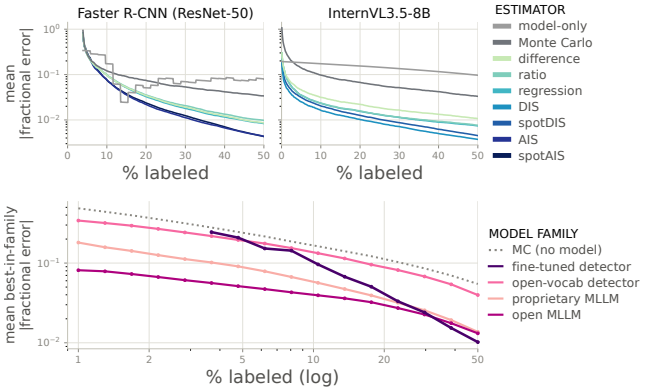

Figure 2: **Top:** mean absolute fractional error over 1,000 trajectories versus labeling budget on *Birds/bird*, one curve per estimator, for fine-tuned Faster R-CNN (left) and frozen InternVL3.5-8B (right). **Bottom:** each family's best case, averaged over 12 tasks; at every task and budget the family is represented by its lowest-error {model, estimator} pair, so a curve is a per-budget optimum for the family, not one fixed model. The fine-tuned detector curve enters at 3.7% labeled due to warmup set.

als are cheap: each unit's prediction is computed once and never changes, so a single cached set of predictions g_0 per model–task pair can be reused across all trials. Fine-tuned models are different, since the model is retrained as labels arrive and its predictions must be recomputed along every trajectory. Running 1,000 such trials for all fine-tuned configurations would have required roughly 82 GPU-years on our compute. We therefore run 5 genuine adaptive trajectories per task–strategy pair and cache the model's predictions at each retraining step (every 25 steps), covering all units in Ω, not only those sampled. Each of the 5 cached sequences is then replayed under 200 different sampling seeds, holding the model's adaptation path fixed while sampling randomness varies, yielding 1,000 trials in total. Hamilton et al. (2026b) used a similar approach with one genuine trajectory. Appendix G provides more details.

5 Results

The top row of Fig. 2 shows estimation error on the *Birds/bird* task against the labeling budget. The ordering of curves previews what the rest of this section quantifies: importance sampling estimators reach the lowest error, uniform sampling estimators follow, and model-only is highest. For the frozen model, the model-only error declines only because revealed units replace predicted counts with true ones (Eq. (1)). Throughout, we put models on an equal footing by scoring each one with its best estimator at every task and budget, and each family by its best model.

Model-assisted counting is generally effective. Model-assisted estimators outperform both the model-only and model-free Monte Carlo baselines on most tasks. We group models into four families: fine-tuned detector, OVD, proprietary MLLM, and open-source MLLM. The left panels of Fig. 3 report each family's error relative to Monte Carlo across occupancy tiers, with values below one indicating improvement. The best model of every family beats Monte Carlo at nearly all budgets, and even the worst model generally does so, except on the full tier.

Open-source MLLMs are the best starter models; fine-tuned detectors overtake with more labels. The best model family changes with the labeling budget. In the bottom panel of Fig. 2, the open-source MLLM curve is lowest from 1% to about 30% labeled. It is already under 10% error at 1% labeled, a budget at which Monte Carlo is still 3.8–6× worse. Fine-tuned detectors start high but improve most rapidly with labels, overtaking every frozen family by about 40%. The bottom-right panel of Fig. 3 shows the same crossover task by task. A model wins a task outright when, scored with its best estimator, its error is strictly the lowest of the 13 models. On *Malaria/difficult* (hatched band), several models predict 0 or 1 on every unit, which the sampling floor maps to the same uniform proposal. So they draw identical units and return identical estimates, and no model is credited. At 5% labeled, open-source MLLMs win 9 of 11 tasks; at 50%, detectors win 8. At the model level, Fig. 4 shows that InternVL3.5-8B wins the most settings overall and at low budgets. Qwen3.7-Plus is the only proprietary MLLM that wins, and OVDs never win. Modern open-source MLLMs are thus a strong training-free starting point, and they substantially outperform earlier-generation OVDs, reflecting progress in vision models, though at a higher inference cost.

Importance sampling estimators are the best default. Labeling where the model predicts the most objects tends to beat correcting a uniform sample. The estimator bars of Fig. 3 show which estimator wins each case, i.e., each {model, task, budget} triple, by budget (top right) and by occupancy tier (beside each panel). Importance sampling estimators win 63% of the 1,560 cases: DIS wins most frequently at every budget, taking 40% of the 1,320 frozen models' cases, and AIS wins 46% of the 240 fine-tuned detectors' cases (the case imbalance makes AIS look minor in the pooled bars). Monte Carlo wins only 14.6% of cases. The tier bars show that DIS's share rises from 17% of the full tier to 45% of the rare tier, uniform sampling estimators win most cases on the full tier but a quarter or fewer on the sparse and rare tiers, and 67% of Monte Carlo's wins are concentrated in the full tier. Together, they suggest that the efficiency of uniform sampling relative to importance sampling declines as the target becomes rare. Still, uniform sampling may be desirable since it makes multi-class estimation (Hamilton et al. 2026a) and data management more straightforward. We therefore suggest DIS or AIS as the default, with uniform sampling as a fallback on dense targets.

Whether to adapt depends on the resources at hand. Off-the-shelf models are a low-barrier entry point through widely available APIs. But proprietary MLLMs can be expensive, roughly $4.60–$9.80 per 1,000 units for the costliest models in our study, and may be unsuitable when data privacy is a concern. Open-source MLLMs can be run locally, and smaller ones cost far less, at roughly $0.25 per 1,000 units. Yet they remain more than 10 times more expensive than Faster R-CNN (Table A6). A fine-tuned detector instead

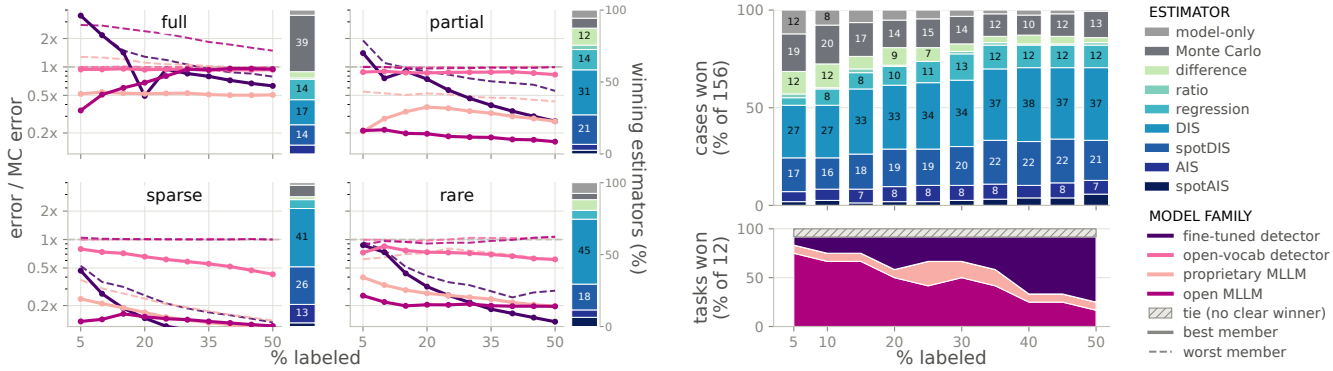

Figure 3: Overview of the main results. A *case* is one {model, task, budget} triple; its winner is the estimator (or, when comparing models, the model) with the lowest error on it. **Left:** error relative to model-free Monte Carlo, one panel per occupancy tier (log scale; below $1\times$ the model helps). Every model is scored with its best estimator; the solid and dashed lines are the family's best and worst model, as medians over the tier's 3 tasks. The bar beside each panel is the share (%) of the tier's 390 cases (3 tasks $\times$ 13 models $\times$ 10 budgets, $t = 5\%, \ldots, 50\%$) won by each estimator. **Top right:** the same shares at each budget t, over the 156 cases (13 models $\times$ 12 tasks) per t. **Bottom right:** share of the 12 tasks whose winning model belongs to each family, at each budget; the hatched band is the one task with no unique winner.

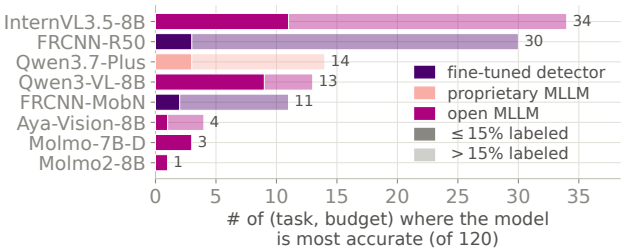

Figure 4: Number of (task, budget) settings, out of 120 (12 tasks $\times$ 10 budgets), on which each model, using its best estimator, has the lowest error of all 13 models. The 10 settings with no unique winner (see text) count for no model. Solid: won at $\leq 15\%$ labeled; faded: won above that. Models that never win are omitted.

pays in labels and training time. One round of retraining the model and recomputing its predictions for all $s \in \Omega$ takes roughly 4 minutes for MobileNet and 8 for the ResNet backbone on average, hours of GPU time over a full run (with a retraining interval of 25 steps), whereas frozen models are scored once. Since the detector's gain also varies across tasks, adaptation is worthwhile when labeling budgets are large, lightweight adaptation is available, or training can be amortized across related deployments, as in our case study (§6). Fully fine-tuning an MLLM requires high compute resources for models with 7B or more parameters. Lightweight adaptation methods and distilling MLLMs into smaller custom detectors are both promising routes for future work.

A small labeled pilot picks a near-optimal frozen model. MLLM performance varies substantially across tasks and models: across the eight MLLMs, the raw predicted totals range from $0.004\times$ to $55\times$ the true count (Fig. A2). Selecting a model before the main loop is therefore worthwhile. We uniformly sample a pilot set $\mathcal{P}$ of 50 units, and rank models by the Hellinger distance between the normalized true and predicted count distributions over $\mathcal{P}$. At 10% and 30% labeled, the model chosen from a single pilot has a median error of only 1.04–$1.06\times$ that of the best candidate, each scored with its best estimator, whether the pool is all 11 frozen models or the eight MLLMs alone. Hellinger distance also recovers the true model ordering better than per-unit RMSE or MAE (Spearman ρ of 0.62–0.70 vs. 0.39–0.50), though for picking a single model the per-unit proxies come close. The weakest proxy is the total-count gap, the absolute difference between a model's predicted and true totals on $\mathcal{P}$ relative to the true total, suggesting how a model distributes predicted mass across units matters more than the accuracy of its magnitude. Appendix H provides further details.

6 Case Study: Pronghorn Aerial Surveys

Estimating how many animals are present in a given landscape is fundamental to robust wildlife management and conservation (Williams et al. 2002). Aircraft make it possible to survey large, complex landscapes, but traditional aerial surveys, in which human observers fly at low altitude, are the most dangerous task that wildlife professionals perform (Johnson, Lindzey, and Guenzel 1991; Sasse 2003). Image-based methods avoid this risk. Without observers on board, aircraft can fly faster and higher, and counts can be controlled for observer biases rather than depending on unrepeatable in-flight observation (Converse et al. 2024). Still, domain shift across landscapes and weather is unavoidable at scale (Converse et al. 2024), so each survey's population estimates must come with a quantified measure of reliability.

We deploy AIS on aerial imagery collected to estimate pronghorn (*Antilocapra americana*) abundance across a 0.575 million-acre region of Wyoming, in collaboration with the Wyoming Game and Fish Department. Pronghorn inhabit open landscapes and are readily visible from above. Where a traditional line-transect survey would have flown observers at $\sim$100 m to sample 10% of the landscape (Johnson, Lindzey, and Guenzel 1991), the same aircraft flew a similar route at a safer 330 m while a 45.7 MP camera captured roughly

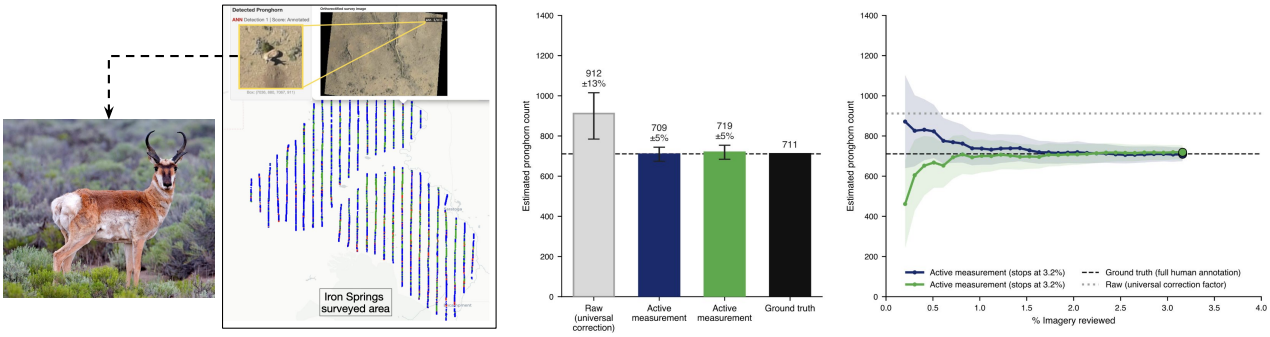

Figure 5: **Left:** a pronghorn (by Alan D. Wilson, CC BY-SA 3.0) and the Iron Springs survey area. Vertical lines are flight transects; each dot is the ground footprint of one survey image (blue: empty unit; green: empty but used for camera pose calibration; red: populated unit). **Center:** estimated pronghorn count from the calibrated model-only baseline (gray) and two AIS runs, with the ground truth. Whiskers are 95% CI. **Right:** the two AIS trajectories and their 95% CIs (shaded) versus % of imagery reviewed. Both reached $\pm 5\%$ uncertainty after reviewing 3.2% of the imagery, the median over all 100 runs (Fig. A7).

one image per second, producing 6,129 images over about 44,000 acres. The collection falls in the *rare* tier, with 98.8% of images containing no pronghorn.

As in the benchmark, we use Faster R-CNN (ResNet-50 FPN, COCO-initialized). Unlike benchmark tasks which start from a small warmup set, a recurring survey can draw on data from earlier surveys. We therefore begin with a detector fine-tuned on four previous surveys from other regions of Wyoming (1,137 annotated tiles containing 1,864 pronghorn), which yields a better g earlier in t. We also use this data to build a better proposal. Each image and its detections are first split into eight non-overlapping tiles. Each tile is described by a fixed-length feature vector of the model's ranked confidence scores, whose length equals the maximum detection count across tiles. A small regression model with a single hidden layer maps this feature vector to a predicted tile count. It is first trained on the earlier surveys and then retrained every 50 steps with the current data, which is far cheaper than fine-tuning the detector. Its predictions are then normalized to form q_t. This ranks predicted-empty tiles more finely than g alone. For comparison, we also evaluate a calibrated model-only baseline that multiplies g by a precision-to-recall ratio estimated from the earlier surveys. We assess its stability by bootstrapping over those four surveys to obtain a distribution of calibration factors.

We ran AIS 100 times with different random seeds (see Fig. A7) and show two representative runs in Fig. 5 (labeled "Active measurement"). Of the 100 runs, 96 reached a $\pm 5\%$ confidence interval (CI) after reviewing only 3.0–3.4% of the imagery. At the median of 3.2%, this is about 1,570 eighth-image tiles, equivalent to roughly 200 full images or one afternoon of work. The other 4 runs stopped after reviewing 27–28.5% of the imagery. Overall, 93 of the 100 CIs contain F_{true}. Both example runs stop at the median and outperform the calibrated model-only baseline. Partial landscape coverage contributes an additional $3\times$ uncertainty, setting a lower bound on the total survey uncertainty. This suggests that once image-level estimates are sufficiently reliable, reducing human effort may matter more than marginal reductions in estimate uncertainty. This case study shows how detector-development costs can be amortized across repeated surveys with lightweight adaptation addressing domain shift, and how each new survey can receive its own reliability estimate without exhaustive annotation.

7 Discussion and Conclusion

We systematically evaluated collection-level counting strategies across 12 tasks spanning biomedical microscopy, wildlife surveys, and overhead imagery of urban structures. Our results show that model-assisted estimation generally outperforms both model-only predictions and model-free Monte Carlo baselines, although the best strategy depends on the labeling budget and task structure. Importance sampling is a strong default when model predictions reliably rank high-count units. MLLMs offer an effective training-free starting point at low budgets, while fine-tuned detectors become increasingly competitive as more labels are acquired. Lightweight adaptation can further adapt models to deployment data without incurring the full cost of detector fine-tuning. Our findings suggest that models for collection-level counting should be selected not solely on the basis of raw counting accuracy, but also on how effectively their predictions support sampling and statistical estimation. Our pronghorn case study demonstrates that these gains translate to practice: AIS achieved a $\pm 5\%$ confidence interval around the true total after reviewing only $\sim$3.2% of the imagery.

Our work also has several limitations. Although the benchmark spans diverse visual domains and occupancy tiers, it is limited to 12 single-class counting tasks. The findings may not fully generalize to multi-class estimation, alternative error metrics, noisy annotations, or settings in which target objects are difficult to verify consistently. Moreover, our comparisons do not capture all deployment costs, including image acquisition, model hosting, annotation interfaces, and the engineering required for repeated adaptation. Finally, our results reflect a snapshot of current frontier models and may need to be revisited as these models continue to improve.

References

Angelopoulos, A. N.; Bates, S.; Fannjiang, C.; Jordan, M. I.; and Zrnic, T. 2023. Prediction-powered inference. *Science*, 382(6671): 669–674.

Bai, S.; Cai, Y.; Chen, R.; Chen, K.; Chen, X.; Cheng, Z.; Deng, L.; Ding, W.; Gao, C.; Ge, C.; et al. 2025. Qwen3-vl technical report. *arXiv preprint arXiv:2511.21631*.

Clark, C.; Zhang, J.; Ma, Z.; Park, J. S.; Tripathi, R.; Lee, S.; Salehi, M.; Ren, J.; Kim, C. D.; Yang, Y.; et al. 2026. Molmo2: Open weights and data for vision-language models with video understanding and grounding. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 28652–28668.

Cochran, W. G. 1977. Sampling techniques.

Converse, R. L.; Lippitt, C. D.; Koneff, M. D.; White, T. P.; Weinstein, B. G.; Gibbons, R.; Stewart, D. R.; Fleishman, A. B.; Butler, M. J.; Sesnie, S. E.; et al. 2024. Remote sensing and machine learning to improve aerial wildlife population surveys. *Frontiers in Conservation Science*, 5: 1416706.

Dash, S.; Nan, Y.; Dang, J.; Ahmadian, A.; Singh, S.; Smith, M.; Venkitesh, B.; Shmyhlo, V.; Aryabumi, V.; Beller-Morales, W.; et al. 2025. Aya vision: Advancing the frontier of multilingual multimodality. *arXiv preprint arXiv:2505.08751*.

Deitke, M.; Clark, C.; Lee, S.; Tripathi, R.; Yang, Y.; Park, J. S.; Salehi, M.; Muennighoff, N.; Lo, K.; Soldaini, L.; et al. 2025. Molmo and pixmo: Open weights and open data for state-of-the-art vision-language models. In *Proceedings of the Computer Vision and Pattern Recognition Conference*, 91–104.

Eikelboom, J. A. J.; Wind, J.; van de Ven, E.; Kenana, L. M.; Schroder, B.; de Knegt, H. J.; van Langevelde, F.; and Prins, H. H. T. 2019. Improving the precision and accuracy of animal population estimates with aerial image object detection. *Methods in Ecology and Evolution*, 10(11): 1875–1887.

Farquhar, S.; Gal, Y.; and Rainforth, T. 2021. On statistical bias in active learning: How and when to fix it. *arXiv preprint arXiv:2101.11665*.

Google DeepMind. 2026. Gemini 3.1 Pro Model Card. https://deepmind.google/models/model-cards/gemini-3-1-pro/. (visited on 2026-07-06).

Gupta, R.; Goodman, B.; Patel, N.; Hosfelt, R.; Sajeev, S.; Heim, E.; Doshi, J.; Lucas, K.; Choset, H.; and Gaston, M. 2019. Creating xBD: A Dataset for Assessing Building Damage from Satellite Imagery. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR) Workshops*.

Hamilton, M.; Lai, J.; Sheldon, D.; and Maji, S. 2026a. Scalable Model-Assisted Multi-Target Estimation in Large Image Collections. arXiv:2607.17581.

Hamilton, M.; Lai, J.; Zhao, W.; Maji, S.; and Sheldon, D. 2026b. Active measurement: Efficient estimation at scale. *Advances in Neural Information Processing Systems*, 38: 145769–145807.

He, K.; Zhang, X.; Ren, S.; and Sun, J. 2016. Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, 770–778.

Howard, A.; Sandler, M.; Chu, G.; Chen, L.-C.; Chen, B.; Tan, M.; Wang, W.; Zhu, Y.; Pang, R.; Vasudevan, V.; et al. 2019. Searching for mobilenetv3. In *Proceedings of the IEEE/CVF international conference on computer vision*, 1314–1324.

Johnson, B. K.; Lindzey, F. G.; and Guenzel, R. J. 1991. Use of aerial line transect surveys to estimate pronghorn populations in Wyoming. *Wildlife Society Bulletin (1973-2006)*, 19(3): 315–321.

Kossen, J.; Farquhar, S.; Gal, Y.; and Rainforth, T. 2022. Active surrogate estimators: An active learning approach to label-efficient model evaluation. *Advances in Neural Information Processing Systems*, 35: 24557–24570.

Lawrence, L.; Saha, O.; Wei, M.; Sun, C.; Maji, S.; and Van Horn, G. 2026. You May Speak Freely: Improving the Fine-Grained Visual Recognition Capabilities of Multimodal Large Language Models with Answer Extraction. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, 1428–1437.

Lempitsky, V.; and Zisserman, A. 2010. Learning To Count Objects in Images. In Lafferty, J.; Williams, C.; Shawe-Taylor, J.; Zemel, R.; and Culotta, A., eds., *Advances in Neural Information Processing Systems*, volume 23. Curran Associates, Inc.

Li, J.; Zhang, X.; Zou, H.; Guo, Y.; Xu, R.; Liu, Y.; Zhu, C.; He, Y.; and Cui, P. 2025. Counts: benchmarking object detectors and multimodal large language models under distribution shifts. In *Proceedings of the Computer Vision and Pattern Recognition Conference*, 9186–9198.

Liu, S.; Zeng, Z.; Ren, T.; Li, F.; Zhang, H.; Yang, J.; Jiang, Q.; Li, C.; Yang, J.; Su, H.; et al. 2024. Grounding dino: Marrying dino with grounded pre-training for open-set object detection. In *European conference on computer vision*, 38–55. Springer.

Ljosa, V.; Sokolnicki, K. L.; and Carpenter, A. E. 2012. Annotated high-throughput microscopy image sets for validation. *Nature methods*, 9(7): 637.

Meng, C.; Liu, E.; Neiswanger, W.; Song, J.; Burke, M.; Lobell, D.; and Ermon, S. 2022. Is-count: Large-scale object counting from satellite images with covariate-based importance sampling. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 36, 12034–12042.

Minderer, M.; Gritsenko, A.; and Houlsby, N. 2023. Scaling open-vocabulary object detection. *Advances in Neural Information Processing Systems*, 36: 72983–73007.

Norouzzadeh, M. S.; Nguyen, A.; Kosmala, M.; Swanson, A.; Palmer, M. S.; Packer, C.; and Clune, J. 2018. Automatically identifying, counting, and describing wild animals in camera-trap images with deep learning. *Proceedings of the National Academy of Sciences*, 115(25): E5716–E5725.

OpenAI. 2025. Introducing GPT-5.2. https://deploymentsafety.openai.com/gpt-5-2. (visited on 2026-07-06).

Owen, A. B. 2013. *Monte Carlo theory, methods and examples*. https://artowen.su.domains/mc/.

Owen, A. B.; and Zhou, Y. 2020. The square root rule for adaptive importance sampling. *ACM Transactions on Modeling and Computer Simulation (TOMACS)*, 30(2): 1–12.

Paiss, R.; Ephrat, A.; Tov, O.; Zada, S.; Mosseri, I.; Irani, M.; and Dekel, T. 2023. Teaching clip to count to ten. In *Proceedings of the IEEE/CVF international conference on computer vision*, 3170–3180.

Perez, G.; Maji, S.; and Sheldon, D. 2024. DISCount: counting in large image collections with detector-based importance sampling. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, 22294–22302.

Qwen. 2026. Qwen3.7-Plus. https://www.qwencloud.com/models/qwen3.7-plus. (visited on 2026-07-06).

Rajaraman, S.; Antani, S. K.; Poostchi, M.; Silamut, K.; Hossain, M. A.; Maude, R. J.; Jaeger, S.; and Thoma, G. R. 2018. Pre-trained convolutional neural networks as feature extractors toward improved malaria parasite detection in thin blood smear images. *PeerJ*, 6.

Ranjan, V.; Sharma, U.; Nguyen, T.; and Hoai, M. 2021. Learning to count everything. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, 3394–3403.

Ren, S.; He, K.; Girshick, R.; and Sun, J. 2015. Faster r-cnn: Towards real-time object detection with region proposal networks. *Advances in neural information processing systems*, 28.

Särndal, C.-E.; Swensson, B.; and Wretman, J. 2003. *Model assisted survey sampling*. Springer Science & Business Media.

Sasse, D. B. 2003. Job-related mortality of wildlife workers in the United States, 1937-2000. *Wildlife society bulletin*, 1015–1020.

Siméoni, O.; Vo, H. V.; Seitzer, M.; Baldassarre, F.; Oquab, M.; Jose, C.; Khalidov, V.; Szafraniec, M.; Yi, S.; Ramamonjisoa, M.; Massa, F.; Haziza, D.; Wehrstedt, L.; Wang, J.; Darcet, T.; Moutakanni, T.; Sentana, L.; Roberts, C.; Vedaldi, A.; Tolan, J.; Brandt, J.; Couprie, C.; Mairal, J.; Jégou, H.; Labatut, P.; and Bojanowski, P. 2025. DINOv3. arXiv:2508.10104.

Wang, W.; Gao, Z.; Gu, L.; Pu, H.; Cui, L.; Wei, X.; Liu, Z.; Jing, L.; Ye, S.; Shao, J.; et al. 2025. Internvl3. 5: Advancing open-source multimodal models in versatility, reasoning, and efficiency. *arXiv preprint arXiv:2508.18265*.

Williams, B. K.; Nichols, J. D.; Conroy, M. J.; and Conroy, M. J. 2002. *Analysis and management of animal populations*. Academic press.

Xia, G.-S.; Bai, X.; Ding, J.; Zhu, Z.; Belongie, S.; Luo, J.; Datcu, M.; Pelillo, M.; and Zhang, L. 2018. DOTA: A Large-Scale Dataset for Object Detection in Aerial Images. In *The IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*.

Xiao, B.; Wu, H.; Xu, W.; Dai, X.; Hu, H.; Lu, Y.; Zeng, M.; Liu, C.; and Yuan, L. 2024. Florence-2: Advancing a unified representation for a variety of vision tasks. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 4818–4829.

Zhang, Y.; Zhou, D.; Chen, S.; Gao, S.; and Ma, Y. 2016. Single-Image Crowd Counting via Multi-Column Convolutional Neural Network. In *2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 589–597.

A Bias and Variance of Estimators

This section shows that every estimator except the model-only baseline is unbiased. The main text evaluates the estimates by their fractional error, but these estimators also yield confidence intervals. Here we write the variance estimators that construct them and report how often the resulting intervals cover F_{true}. We use the notation of §3.1, and let $\mathcal{H}_t$ denote the history at step t: the revealed units, their labels, and the fitted g_t, c_t, q_t, and h_t. Given $\mathcal{H}_t$, the only randomness at step t is the draw $s_t \sim q_t$ from $\mathcal{R}_t$.

Unbiasedness of the incremental estimators Every incremental estimator of §3.3 belongs to one family, written with the per-unit residual $d_t(s) = f(s) - c_t g_t(s)$ as

$$\hat{F}_t = F(\mathcal{D}_t) + c_t\, G_t(\mathcal{R}_t) + \frac{d_t(s_t)}{q_t(s_t)}. \tag{A1}$$

The incremental difference, ratio, and regression estimators are the cases $c_t \in \{1, \hat{\gamma}_t, \hat{\beta}_t\}$ under uniform sampling, where $q_t \equiv 1/M_t$ and the last term becomes $M_t\, d_t(s_t)$, recovering Eq. (3). AIS, DIS, spotAIS, and spotDIS are the case $c_t = 0$, recovering Eq. (4).

Given $\mathcal{H}_t$, the first two terms of Eq. (A1) are constants, and the last term is an importance-weighted estimate of the residual total on $\mathcal{R}_t$. With $q_t(s) > 0$ on $\mathcal{R}_t$,

$$\begin{aligned}
&\mathbb{E}\big[\hat{F}_t \mid \mathcal{H}_t\big] \\
&\quad = F(\mathcal{D}_t) + c_t G_t(\mathcal{R}_t) + \sum_{s \in \mathcal{R}_t} q_t(s)\, \frac{d_t(s)}{q_t(s)} \\
&\quad = F(\mathcal{D}_t) + c_t G_t(\mathcal{R}_t) + F(\mathcal{R}_t) - c_t G_t(\mathcal{R}_t) \\
&\quad = F(\mathcal{D}_t) + F(\mathcal{R}_t) \;=\; F_{\text{true}}.
\end{aligned}$$

The last term removes, in expectation, exactly the error of the scaled prediction $c_t G_t(\mathcal{R}_t)$, whatever g_t and c_t are. This conditional expectation is the same for every history, and $\hat{F}_t$ is bounded because Ω is finite, $q_t > 0$ on $\mathcal{R}_t$, and c_t is finite. So the law of total expectation gives $\mathbb{E}[\hat{F}_t] = F_{\text{true}}$. We only need that q_t is positive on $\mathcal{R}_t$ and that g_t, c_t, and q_t are determined by $\mathcal{H}_t$, so unbiasedness holds however the model and proposal adapt; the accuracy of g_t affects only the variance.

The fallback to $c_t = 1$ (§3.3) applies while fewer than 10 units have been revealed, and also when $G_t(\mathcal{D}_t) < 1$ for the ratio coefficient or when g_t is constant on $\mathcal{D}_t$ for the regression coefficient. It depends only on $\mathcal{H}_t$, so it keeps c_t finite without affecting unbiasedness.

The Monte Carlo baseline (Eq. (2)) is not a member of this family. It is the classical expansion estimator applied to the $N - w$ non-warmup units, where w matches the model it is compared with ($w = 50$ for fine-tuned detectors and $w = 0$ for frozen models), and is unbiased under uniform sampling without replacement (Cochran 1977).

The COMB and LURE weights (Appendix C) are positive, sum to one, and depend only on N, w, and the step index rather than on any observed label. Thus their weighted sum is also unbiased:

$$\mathbb{E}\big[\hat{F}_{1:t}\big] = \sum_{i=1}^{t} \bar{\alpha}_{i,t}\, \mathbb{E}\big[\hat{F}_i\big] = F_{\text{true}}.$$

Variance and confidence intervals For all incremental estimators, we use the "simple" conditional variance estimator of Hamilton et al. (2026b),

$$\widehat{\mathrm{Var}}^{\,\mathrm{simp}}_{1:t} = \sum_{i=1}^{t} \bar{\alpha}^2_{i,t} \big(\hat{F}_i - \hat{F}_{1:t}\big)^2. \tag{A2}$$

They also define an alternative that estimates the conditional variance separately at each step, and report that the simple form behaves better where coverage is poor. Confidence intervals are $\hat{F}_{1:t} \pm z_{1-\alpha/2}(\widehat{\mathrm{Var}}^{\,\mathrm{simp}}_{1:t})^{1/2}$, where $z_{1-\alpha/2}$ is the $1 - \alpha/2$ quantile of the standard normal distribution.

The two batch estimators are handled separately. The Monte Carlo baseline takes the classical without-replacement variance with a finite-population correction (Cochran 1977, Ch. 2), applied to the non-warmup units because the warmup is counted exactly:

$$\widehat{\mathrm{Var}}^{\,\mathrm{MC}}_{1:t} = \frac{K\,(K-t)}{t}\,\hat{\sigma}^2_t,$$

where $K = N - w$ and $\hat{\sigma}^2_t$ is the sample variance of the t observed counts. The model-only baseline has no variance estimator or confidence interval: given the model snapshot and the revealed set, it is deterministic and has no correction term from which a sampling error could be estimated.

Confidence interval coverage Table A1 reports how often the nominal 95% interval from a single trial contains F_{true}, and how wide it is. Excluding the spot-checked variants, coverage is below nominal but improves with budget (Table A2): it averages 0.89 over the ten budgets and rises from 0.82 at the 5% budget to 0.92 at 50%, consistent with Hamilton et al. (2026b). Among these estimators, AIS undercovers most at the smallest budgets (0.62 at 5%), but is in line with the other estimators from 20% onward.

Coverage of the spot-checked variants On the sparse and rare tiers, the spot-checked variants have the lowest coverage in Table A1 (0.57–0.79). Here, most units are empty and predicted to be empty, so q_t^{spot} spreads its $\lambda = 5\%$ of the probability mass over many units (Appendix B). A populated unit wrongly routed there therefore receives a very low sampling probability and hence a very large importance weight if drawn. Most trials never draw such a unit and underestimate both the total and its variance; the few that do overestimate both. The two effects cancel on average, so the mean estimate over trials remains close to F_{true} (mean relative error -0.001), but a single trial usually reports an interval too narrow to reach F_{true}. Consistent with this, each spot-checked variant reports a narrower interval than its non-spot counterpart on these tiers. This explanation is specific to the sparse and rare tiers. It does not account for the below-nominal coverage of AIS and spotAIS on the full and partial tiers.

Relevance to the case study These coverage results do not transfer to the pronghorn case study of §6. It is a rare-tier collection, but its detector was fine-tuned on earlier surveys of the same species rather than starting from a small warmup, and its proposal is specifically adapted for the task to handle domain shift between deployments. Its coverage is therefore

	Coverage				*Half-width / F_{true}*			
Estimator	full	part	spar	rare	full	part	spar	rare
Monte Carlo	0.94	0.94	0.93	0.86	0.07	0.15	0.20	0.49
difference	0.89	0.93	0.88	0.88	0.09	0.12	0.18	0.47
ratio	0.87	0.92	0.88	0.86	0.12	0.12	0.16	0.43
regression	0.90	0.92	0.88	0.86	0.07	0.11	0.16	0.41
DIS	0.92	0.92	0.93	0.86	0.08	0.11	0.09	0.19
AIS	0.83	0.86	0.92	0.87	0.06	0.07	0.04	0.12
spotDIS	0.91	0.89	0.79	0.61	0.10	0.10	0.06	0.09
spotAIS	0.79	0.80	0.71	0.57	0.08	0.07	0.03	0.04
excl. spot	0.89	0.92	0.89	0.86	0.09	0.11	0.15	0.38

Table A1: Behavior of 95% confidence intervals by occupancy tier. *Coverage* is the fraction of the 1,000 trials of a {model, task, budget} case (a {task, budget} case for Monte Carlo, which uses no model) whose interval contains F_{true}, averaged over the cases in a tier. *Half-width* is $1.96\,(\widehat{\text{Var}}^{\text{simp}}_{1:t})^{1/2}/F_{\text{true}}$ (with $\widehat{\text{Var}}^{\text{MC}}_{1:t}$ for Monte Carlo), summarized by its median within a case and again across cases. The *excl. spot* row pools all cases of the six estimators other than spotDIS and spotAIS.

	Coverage				*Half-width / F_{true}*			
Estimator	5%	10%	20%	50%	5%	10%	20%	50%
Monte Carlo	0.86	0.90	0.92	0.94	0.38	0.28	0.19	0.10
difference	0.82	0.87	0.90	0.92	0.37	0.27	0.18	0.09
ratio	0.81	0.86	0.89	0.91	0.41	0.31	0.19	0.10
regression	0.81	0.86	0.89	0.92	0.36	0.26	0.17	0.08
DIS	0.87	0.89	0.91	0.92	0.31	0.21	0.14	0.07
AIS	0.62	0.80	0.90	0.92	0.22	0.17	0.12	0.03
spotDIS	0.73	0.76	0.79	0.85	0.20	0.13	0.10	0.05
spotAIS	0.47	0.63	0.74	0.80	0.12	0.08	0.05	0.02
excl. spot	0.82	0.87	0.90	0.92	0.36	0.26	0.17	0.08

Table A2: Behavior of 95% confidence intervals by labeling budget, at four of the ten evaluated budgets (5% to 50% in steps of 5%). All rows are defined and summarized over cases as in Table A1, but pooled across tiers.

measured separately: of 100 AIS runs with different seeds, 93 final 95% intervals contain F_{true} (Fig. A7).

B Spot-Checked Proposal Construction

The spot-checked proposals of §3.3 reserve a small uniform budget for units predicted to be empty. This section describes how the classifier h_t behind that prediction is fit and how its threshold is set. The two variants differ only in g_t: spotDIS keeps $g_t = g_0$, while spotAIS updates g_t as the detector is retrained. In both, h_t is refit on $\mathcal{D}_t$ every 25 steps, with empty units ($f(s) = 0$) as the positive class. Until at least two empty and two populated units have been revealed, h_t is left unfit and flags nothing, so only units with $g_t(s) = 0$ are predicted empty. When no remaining unit is predicted empty, all probability goes to q_t^{count}; when none is predicted populated, all goes to q_t^{spot}. Hence, $q_t(s) > 0$ for every remaining unit, the estimator is unbiased for any h_t, and errors in h_t affect only the variance.

We set h_t's threshold to limit that variance, exploiting an asymmetry between the two kinds of error. A populated unit routed to q_t^{spot} receives a small probability and hence a large importance weight $f(s)/q_t(s)$ when drawn, whereas an empty unit routed to q_t^{count} costs little, since $f(s) = 0$ contributes nothing to the estimate. We therefore place the threshold just above the $(1 - \phi)$ quantile of the scores of the revealed populated units, with $\phi = 0.01$, so that at most a fraction ϕ of them would be flagged as empty. Scores are out-of-fold, from stratified cross-validation over $\mathcal{D}_t$ with five folds (fewer when a class has fewer than five revealed units), because in-sample scores are over-confident. While the number n of revealed populated units is below 100, $\lfloor \phi n \rfloor = 0$ and none of them is flagged.

C Estimator Weight Sequences

The combination weights of §3.3 are $\bar{\alpha}_{i,t} = \alpha_i / \sum_{j=1}^{t} \alpha_j$. When the model is fine-tuned, we use the COMB weights of Hamilton et al. (2026b), the product of the LURE weight for without-replacement sampling (Farquhar, Gal, and Rainforth 2021) and the square-root rule for adaptive proposals (Owen and Zhou 2020):

$$\alpha_i^{\text{COMB}} = \omega_i^{\text{LURE}} \sqrt{i},$$
$$\omega_i^{\text{LURE}} = \frac{1}{M_i\,(M_i - 1)},$$

where $M_i = |\mathcal{R}_i| = N - w - i + 1$. So ω_i^{LURE} grows as the remaining pool shrinks. When the model is not updated, we drop the square-root factor and use $\alpha_i = \omega_i^{\text{LURE}}$. The weights require $M_i \geq 2$, which the budget cap $t \leq N/2$ guarantees for every task.

D Tile Sizing and Pool Construction

We write *image* for a source image before tiling; a *unit* of Ω is either a tile or a whole image.

Tile-size rule This section describes how we choose each task's tile size. Raw object size varies by nearly an order of magnitude across tasks: the median longer side of the sampled objects ranges from 24 px for *Birds/bird* to 190 px for *DOTA/tennis court*. A single tile size would therefore present very different apparent object scales across tasks, confounding the comparison across occupancy tiers. Let m be the median longer side of the sampled boxes described below. We use Faster R-CNN's internal resizing, which resizes each unit to a short side of 800 px (Appendix F), as the reference because it is the only model type we fine-tune. An object of size m in a tile of side ℓ then appears at $800\,m/\ell$ px, so the 80 px target gives

$$\ell \;=\; \text{round}_{16}\big(\tfrac{800}{80}\,m\big) \;=\; \text{round}_{16}(10\,m), \qquad \text{(A3)}$$

where round_{16} rounds to the nearest multiple of 16 to match detector strides. We chose 80 px to stay well above 32 px, which is both the smallest anchor size of the detectors and the upper limit of COCO's "small" objects (area below 32^2 px^2).

An image whose longer side is at most 1.5ℓ is kept whole, since tiling it would produce one full tile plus thin strips. The size check is applied per image, so a task may contain both whole-image and tile units. A larger image of width W and height H yields a $\lceil W/\ell \rceil \times \lceil H/\ell \rceil$ grid of tiles, which are square except at the right and bottom edges.

Object samples for tile-size estimation For each task, images are visited in a seeded random order, and at most three randomly chosen bounding boxes are read per image until 30 boxes from at least 10 images are collected. *Birds* has only two images, so all of its boxes come from those two. Only each box's longer side enters Eq. (A3), so the sample reveals no unit count, and the sampled images stay in the pool. As a diagnostic, we check that the apparent sizes $800\,x/\ell$ of the sampled longer sides x have first and third quartiles between 32 and 150 px. Among the 12 tasks, only *DOTA/plane* falls outside, with an upper quartile at 214 px; we keep the task and its tile size as is.

Pool construction Tiling every available image would make some tasks too large to run the benchmark efficiently. For example *DOTA/small vehicle* reaches 5,032 units from only 64 of its 1,179 images. We therefore cap each task at roughly 5,000 units. Images are admitted whole, in a seeded random order, until the running unit count first reaches the cap. An image's unit count is computed from its dimensions alone, so admission does not depend on object counts. Ω and F_{true} are defined on the admitted images, fixed before any model is run, and shared by all methods. The images sampled for tile sizing are taken from the front of the same order.

E Class Selection

Table A3 lists every class in the four datasets; none was filtered out before tiering. Each row is measured on that class's own pool Ω (Appendix D), so the empty fraction is a property of the (class, tile size) pair rather than of the class alone. Because tile size normalizes apparent object scale, the tiers compare classes at roughly matched object size rather than matched tile area. Within each tier we rank by F_{true}, since a larger total makes the fractional error less sensitive to single objects: one object is 0.6% of *DOTA/container crane*'s total of 156, but only 0.001% of *Malaria/red blood cell*'s 83,034. We take the three largest classes in the full, partial, and sparse tiers. In the rare tier, the three largest are giraffe, helicopter, and elephant. We skip elephant and the fourth-largest, zebra, and take storage tank instead, so that the tier does not contain two similar aerial wildlife tasks. Storage tank also has the most dispersed count distribution in the benchmark (CV = 14.2).

F Detector Training Details

This section gives the details needed to reproduce the two fine-tuned Faster R-CNN detectors.

Warmup set The warmup set $\mathcal{W}$ is drawn once per trajectory using ground-truth occupancy: 25 units are sampled uniformly from the populated units, and the other 25 uniformly from the rest of Ω, without replacement. This guarantees that at least half of $\mathcal{W}$ is populated. Because the warmup counts toward the labeling budget, on tasks with few units it can exceed the smallest budgets (e.g., 14% of the 347 units of *Malaria/difficult*), so detector results there are reported at the first step after the warmup; the detectors still do not win these tasks at those budgets, so the conclusions are unchanged.

Implementation and inference Both detectors are the torchvision implementations with COCO_V1 weights. Each unit is preprocessed by the detector's default transform, which normalizes by ImageNet channel statistics and resizes the unit to a short side of 800 px with the long side capped at 1,333 px; we add no data augmentation. We raise the cap on detections per unit from the torchvision default of 100 to 300, since five of the twelve tasks have units with more than 100 objects (up to 219). All other inference defaults are unchanged, including non-maximum suppression at IoU 0.5. The predicted count $g_t(s)$ is the number of detections on unit s with confidence at least 0.5.

Training targets A unit's annotations are filtered to the target class, so objects of other classes are unlabeled and act as negatives. Each retraining step uses every unit revealed so far, including empty ones. While counts follow the center-point rule (§4.1), training targets are all boxes overlapping the unit, clipped to its bounds. Under the center-point rule, an object straddling a boundary would be visible but unlabeled in the neighboring unit, teaching the detector to treat it as background. The cost is that a straddling object can be detected in both units, so g_t can overcount there. This error enters the model-only baseline directly, but affects only the variance of the model-assisted estimators.

Fine-tuning We fine-tune the feature pyramid, region-proposal network, RoI heads, and the deepest backbone stage (`layer4` for ResNet-50; the final four inverted-residual blocks for MobileNetV3-Large), and freeze all earlier backbone stages. Each retraining step continues from the previous checkpoint, runs five epochs over the revealed units under the standard Faster R-CNN losses, and then recomputes the predictions for all of Ω. We use SGD with momentum 0.9, weight decay 10^{-4}, and batch size 4, rebuilding the optimizer at each retraining step. Learning rates are constant: 5×10^{-4} for the region-proposal network and RoI heads, and 5×10^{-5} for the feature pyramid and the unfrozen backbone stage.

Computing infrastructure and cost Runs used one GPU per job (at least 16 GB of VRAM) on a heterogeneous cluster, under PyTorch 2.7.1 and torchvision 0.22.1 built against CUDA 11.8. The benchmark comprises 360 genuine detector trajectories: 2 backbones $\times$ 12 tasks $\times$ 3 samplers (uniform, AIS, spotAIS) $\times$ 5 seeds. The seed fixes the warmup draw, the sampler, and the training shuffles; we do not enable deterministic GPU kernels, so a rerun reproduces the protocol but not bitwise-identical weights. In job wall-clock time, the 360 trajectories cost about 3,600 GPU-hours (804 for MobileNetV3-Large and 2,806 for ResNet-50), including dataset construction, the initial fit on $\mathcal{W}$, and the overhead of resuming preempted jobs. Scaling each of the 72 settings from 5 to 1,000 genuine trajectories at its measured mean cost gives the 82 GPU-year estimate of §4.2.

Tier	Dataset	Class	%∅	μ	CV	N	F_{true}	img. pool	img. total	tile px
full ($< 5\%$ empty)	Malaria	**red blood cell**	2.1	49.19	0.77	1,688	83,034	1,328	1,328	1,168
	Malaria	**trophozoite**	0.0	2.42	0.72	655	1,584	655	655	1,488
	Malaria	**difficult**	0.0	1.29	0.44	347	446	347	347	1,424
	Malaria	schizont	0.0	1.14	0.35	166	190	166	166	1,536
	Malaria	gametocyte	0.0	1.06	0.23	147	156	147	147	1,408
	Malaria	leukocyte	0.0	1.10	0.27	94	103	94	94	1,248
partial (5–60% empty)	Birds	**bird**	24.4	14.36	1.48	1,302	18,696	2	2	240
	DOTA	**harbor**	52.1	3.90	2.40	2,084	8,118	456	456	1,360
	DOTA	**tennis court**	37.5	4.63	1.19	689	3,188	401	401	1,888
	Malaria	ring	30.1	1.11	0.96	472	522	301	301	1,280
	DOTA	soccer ball field	47.9	1.08	1.75	449	487	189	189	2,336
	DOTA	ground track field	56.0	0.68	1.96	696	476	251	251	2,384
sparse (60–90% empty)	DOTA	**ship**	78.7	2.10	5.07	5,061	10,632	160	573	512
	DOTA	**small vehicle**	80.8	1.31	5.37	5,032	6,583	64	1,179	256
	DOTA	**plane**	79.4	0.67	3.21	5,043	3,370	90	269	544
	DOTA	large vehicle	85.4	0.54	4.28	5,006	2,682	73	703	304
	DOTA	swimming pool	88.7	0.25	4.07	5,022	1,274	131	288	432
	DOTA	bridge	85.1	0.23	2.87	5,041	1,136	125	288	608
	DOTA	basketball court	78.6	0.64	2.76	1,046	672	155	155	1,152
	DOTA	baseball diamond	79.4	0.43	2.51	1,439	625	174	174	1,152
	DOTA	container crane	80.9	1.03	3.68	152	156	9	9	992
rare ($\geq 90\%$ empty)	Wildlife	**giraffe**	91.7	0.16	4.39	5,005	806	143	209	752
	DOTA	**helicopter**	92.8	0.47	8.12	1,523	713	48	48	768
	Wildlife	elephant	94.9	0.12	6.11	5,000	608	97	195	624
	Wildlife	zebra	95.9	0.12	7.19	5,112	595	42	189	384
	DOTA	**storage tank**	98.1	0.11	14.24	5,176	549	26	262	272
	DOTA	roundabout	97.3	0.03	6.15	5,064	145	58	249	384

Table A3: All candidate classes, grouped by occupancy tier and ordered by decreasing pool total F_{true} within each tier. Selected benchmark classes are shown in **bold**. μ is the average target-object count per unit and *CV* its coefficient of variation across units (dispersion); N is the total number of estimation units; F_{true} is its exact target-object count within the N units; *img. pool* and *img. total* are the numbers of included and available source images, respectively; and *tile px* is the class-specific tile side length. Pools may slightly exceed the $5{,}000$-unit target because the final admitted image is retained in full.

G Trajectory Simulation Details

Because genuine fine-tuned trajectories are expensive in time, most fine-tuned trials are replays (§4.2). This section describes replay and tests whether it changes the results.

Replay A genuine fine-tuned trajectory caches the predictions g_t on all of Ω at every retraining step and, for spotAIS, the classifier h_t. In a genuine trajectory, the revealed labels determine how the model retrains, which in turn affects later sampling. A replay severs this feedback loop: it reuses the original warmup set and cached snapshots and varies only the sampling randomness. A replay may therefore use snapshots trained on labels it did not itself reveal, in the proposal for AIS and spotAIS, in the correction terms for the model-assisted uniform sampling estimators, and directly in the model-only estimate. Frozen-model trials are all genuine: g_0 is computed once and reused, and for spotDIS the classifier h_t is cheap enough to refit within every trial. The Monte Carlo baseline uses no model at all. It estimates the total by scaling up the average of the true counts it has seen, so its trials are also cheap to run genuinely. It also provides an exact control for replay. Computed on a replay of a fine-tuned path, Monte Carlo ignores the cached snapshots, so the replay differs from a genuine trial only in its uniform draws, which follow the same distribution. For each fine-tuned setting, one of the 200 sampling seeds per cached model path reproduces the genuine trajectory and the other 199 are replays, so the $1{,}000$ trials comprise five genuine trajectories and 995 replays.

Validation metrics We use two metrics, both computed from $\hat{F}_{1:t}$ at the 10, 20, 30, 40, and 50% budgets. Each has a known value under faithful replay, that is, when genuine trajectories and replays are exchangeable.

Held-out coverage asks whether a genuine estimate looks like a replay of a different path. For each genuine trajectory, we take the $4 \times 199 = 796$ replays of the other four paths of the same task and check whether the genuine estimate falls inside their central 90% interval. We report the fraction of the 60 genuine trajectories (twelve tasks, five paths) for which it does, averaged over the five budgets. Under faithful replay, this fraction should be about 0.90.

Landing shift asks whether a genuine estimate sits systematically above or below the replays of its own path, which removes between-path variation. Let X_r be the estimate of genuine trajectory r, and μ_r and σ_r the mean and standard deviation of its 199 replays. The landing position $u_r = (X_r - \mu_r)/\sigma_r$ places the genuine estimate among its replays, in replay standard deviations. Under faithful replay,

Estimator	MobileNetV3		ResNet-50	
	Cov.	Shift	Cov.	Shift
faithful replay, in expectation	*.90*	*0*	*.90*	*0*
model-only	.61	-1.47	.64	-0.96
Monte Carlo *(exact, measured)*	.91	$+0.10$	.91	$+0.10$
difference	.86	$+0.23$	.89	$+0.04$
ratio	.86	$+0.17$	.88	$+0.10$
regression	.88	$+0.21$	.88	$+0.02$
AIS	.90	-0.01	.86	-0.06
spotAIS	.92	-0.03	.86	$+0.05$

Table A4: Replay validation for the two fine-tuned detectors. *Cov.* is held-out coverage and *Shift* is landing shift, both defined in Appendix G. The first row gives their values under faithful replay; shifts within ± 0.28 are consistent with it. Monte Carlo ignores the detector, so its entries are the same under both.

u_r has mean zero and a standard deviation close to one (1.008 if the estimates are normally distributed); we observed that the exact Monte Carlo control gives 0.97. The landing shift is the average of u_r over the five budgets, then over the five trajectories of each task, and then over the twelve tasks, so it is zero in expectation under faithful replay.

Averaging makes the landing shift far less variable than a single u_r, so we judge it by its standard error rather than by the standard deviation of u_r. Because trajectories from the same task may share a discrepancy, we compute the standard error from the twelve task means rather than the 60 trajectories; it is approximately 0.14. Under faithful replay, the landing shift would then fall within two standard errors of zero, ± 0.28, roughly 95% of the time, so we treat shifts in this range as consistent with faithful replay.

Fidelity for the model-assisted estimators Table A4 reports both metrics, and Fig. A1 shows the landing positions at the 30% budget. For the model-assisted estimators, replay is close to genuine. The Monte Carlo control gives a held-out coverage of 0.91 and a shift of $+0.10$, illustrating the variation that remains with only five paths per task even when replay is exact. Across the ten pairs of a model-assisted estimator (AIS, spotAIS, difference, ratio, or regression) and a detector, mean held-out coverage is 0.88, and landing shifts range from -0.06 to $+0.23$, all inside the ± 0.28 range expected under faithful replay. The largest shifts, $+0.23$ and $+0.21$ for difference and regression with MobileNetV3, are approximately 1.6 and 1.5 standard errors above zero. These results do not establish that replay is equivalent to running $1{,}000$ genuine trajectories. They indicate, however, that any discrepancy is small for these estimators: 0.03 in mean held-out coverage relative to the exact control, and at most 0.23 replay standard deviations in landing shift.

Fidelity for the model-only baseline Replay is less faithful for the model-only baseline: its held-out coverage is only 0.62 on average, and its shifts are -1.47 and -0.96 for MobileNetV3 and ResNet-50, the only table entries clearly outside sampling variation (see also the top row of Fig. A1). Genuine trajectories thus tend to produce lower estimates than their replays. For MobileNetV3, the shift grows in magnitude from -0.5 at the 10% budget to -2.1 at 50%, consistent with increasing divergence between a replay's revealed set and the set used to train its cached snapshot. The model-only baseline has no correction term, so this mismatch enters the estimate directly. At the 50% budget, genuine trajectories have 9.4% mean absolute fractional error, compared with 10.5% for the replays, so our reported results may slightly overestimate the error of the model-only baseline.

Limitations The analysis uses only five genuine trajectories per setting, which limits our estimate of between-path variability, especially at low budgets. The coefficient of variation of the predicted pool total $G_t(\Omega)$ across the five paths exceeds 0.9 on nine of the twelve tasks below the 20% budget, and its median falls to 0.04 from 30% onward. Replay estimates of the low-budget spread are therefore less certain than estimates of its center.

H Pilot Study Details

The eight MLLMs disagree widely: within a task their predicted totals can differ by orders of magnitude, and each model's accuracy shifts from task to task (Fig. A2). A model must therefore be chosen per task, and a small labeled pilot can do so before the main loop. This section describes the pilot procedure, the four proxies we compare, and how we evaluate them. The study covers only the 11 frozen models; a one-shot pilot cannot represent a detector that retrains and improves as labels accumulate. For each pilot, we draw 50 units uniformly at random from Ω, reveal their true counts, and score every candidate model on the same units. Because a pilot is itself a random draw, we repeat the procedure over 300 independent pilots, on two candidate sets (all 11 frozen models, and the 8 MLLMs alone) and at two labeling budgets (10% and 30%).

The four proxies Let f_i and g_i be the true and predicted counts on pilot unit i. The Hellinger distance compares $p_i = f_i / \sum_j f_j$ with $q_i = g_i / \sum_j g_j$,

$$H(p, q) = \tfrac{1}{\sqrt{2}} \Big(\sum_i (\sqrt{p_i} - \sqrt{q_i})^2 \Big)^{1/2},$$

so it depends only on how the predicted mass is spread across units and not on its magnitude. The other three proxies are scale-sensitive: the per-unit RMSE $\big(\frac{1}{|\mathcal{P}|} \sum_i (g_i - f_i)^2 \big)^{1/2}$ and MAE $\frac{1}{|\mathcal{P}|} \sum_i |g_i - f_i|$, and the total-count gap $|\sum_i g_i - \sum_i f_i| / \sum_i f_i$. The Hellinger distance and the total-count gap are undefined on a pilot that contains no objects (the former also when a model predicts zero on every pilot unit). A model whose proxy is undefined is excluded from that pilot's selection, and ρ averages each proxy over the pilots on which it is defined.

Proxy evaluation Table A5 evaluates each proxy with three metrics, all of which compare proxy values with the models' realized errors. A model's realized error is the lowest mean absolute fractional error it reaches at that budget in our benchmark runs, over the seven estimators available to a

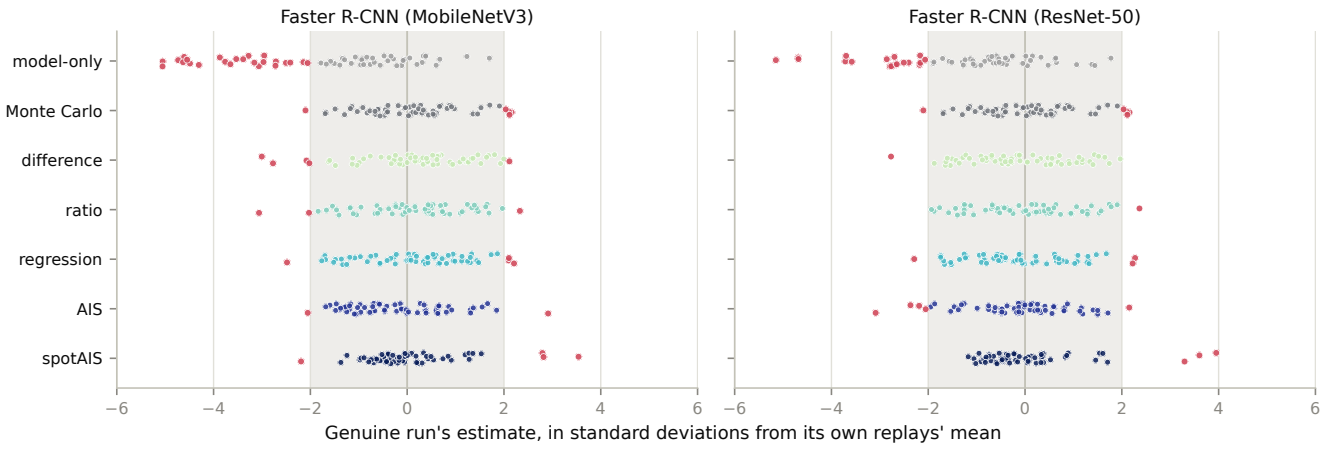

Figure A1: Landing positions u_r at the 30% budget, underlying the *Shift* column of Table A4. Each point is one genuine trajectory, placed relative to the 199 replays of its own cached model path in replay standard deviations (0 is the replay mean). The shaded band spans ± 2 standard deviations for an individual trajectory, not the ± 0.28 interval for the mean shift; points outside it are red. The model-assisted estimators resemble the exact Monte Carlo control, whereas the model-only baseline has a pronounced negative tail.

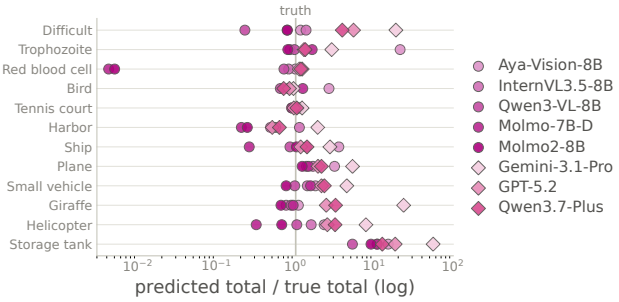

Figure A2: Total count predicted by each of the eight MLLMs as a ratio to the true total, on every task (rows, ordered full to rare; log axis, 1 = exact); circles are open-source MLLMs and diamonds are proprietary ones, each tinted by model.

frozen model. The first metric, ρ, measures how well a proxy orders all candidate models. It is the within-task Spearman rank correlation between the models' proxy values, averaged over the 300 pilots, and their realized errors. Because averaging removes pilot noise, ρ can be interpreted as an optimistic estimate of how well a single pilot orders the candidates. The other two metrics measure how good a single 50-unit pilot is at selecting the best model. *beats* is the fraction of the other candidates whose error is higher than that of the pick (mean over pilots), and *ratio* is the pick's error divided by that of the actual best model among candidates (median over pilots). For all three metrics, the ideal value is 1.0.

Proxy comparison The Hellinger distance is the most reliable proxy, and the total-count gap is the weakest on every measure. For selecting a single model, the per-unit proxies (RMSE and MAE) come close to the Hellinger distance on the MLLM-only set (*ratio* of 1.06–1.15× against 1.04–1.06×). But they trail Hellinger on the MLLM + OVD set (1.12–1.22× against 1.06×). The Hellinger distance's clearer advantage is in ρ (0.62–0.70, against 0.39–0.50 for RMSE and MAE and 0.25–0.38 for the total-count gap): it orders all candidates rather than only identifying the best one, which matters when a practitioner wants to weigh several candidates. For choosing a frozen model, we therefore recommend ranking candidates by the Hellinger distance on the pilot. It rewards models whose predictions spread across units like the true counts, rather than the accuracy of their magnitude.

Full-pool comparison Figs. A3 to A6 compare the total-count gap and the Hellinger distance as predictors of downstream performance when both are computed on all of Ω rather than on a pilot, for each frozen model on each task. Downstream performance is measured for DIS by the mean AUC, the area under the curve of fractional error against the percentage of units labeled (lower is better), and by the fractional error at 10% labeled. Against both measures, the Hellinger distance has a higher R^2 than the total-count gap $|G_0(\Omega) - F_{\text{true}}|/F_{\text{true}}$, so it explains more of the variation in downstream performance. For example, on *Birds/bird*, Molmo-7B-D-0924 has a smaller total-count gap than Aya Vision-8B but a much larger Hellinger distance, and a higher AUC and error at 10% labeled.

I Model Inference, Cost, and Prompting

Here we compare the inference cost of all 13 models, and describe how the MLLMs were run and how counts were extracted from their responses. Table A6 converts each model's cost to dollars per 1,000 units, so that locally run and API-based models can be compared on one scale. Table A7 lists the raw measurements behind its MLLM rows (GPU-hours per unit or API cost per unit), along with each MLLM's source and count-extraction method.

MLLM inference and cost estimation The open-source models (Qwen3-VL-8B, InternVL3.5-8B, Molmo-7B-D-0924, Aya Vision-8B, Molmo2-8B) were each run on a single A100 GPU over all units of the benchmark (Table A7). Wall-clock time was measured from job start to the last prediction written, and totals approximately 15.4 GPU-hours across the five models ($2.67 + 4.30 + 2.69 + 2.89 + 2.87$). To measure latency and cost, the proprietary models (GPT-5.2, Gemini 3.1 Pro, Qwen3.7-Plus) were queried via the OpenRouter API on a sample of 60 units (5 per task; 59 for Qwen3.7-Plus due to one API failure), with cost computed from the recorded token usage and per-token pricing. The samples cost \$0.2787 (GPT-5.2), \$0.5861 (Gemini 3.1 Pro), and \$0.3488 (Qwen3.7-Plus). Extrapolated to the full benchmark, the estimated costs are \$156.58, \$329.35, and \$199.33, with estimated wall times of 60.2, 67.3, and 759.6 hours.

MLLM count extraction All five prompts are collected in Table A8, one per row. For open-source models, the *open-source count* prompt is followed by a second in-model pass (*open-source count extraction*) that re-prompts the model to map its free-form response to a single integer, producing a structured choice (`nlg2choice`) parsed via `int(nlg2choice)` with a regex-digit fallback. For proprietary models, the *proprietary count* prompt asks the model to reason briefly and finish with a `Final answer: <number>` line. A separate extractor call (GPT-5.2 via OpenRouter) runs the *proprietary count extraction* prompt, which uses the number on the final-answer line when present and falls back to a `final_answer_number` regex over the raw response. An optional *verification* pass checks whether the extracted count is plainly present in the original response.

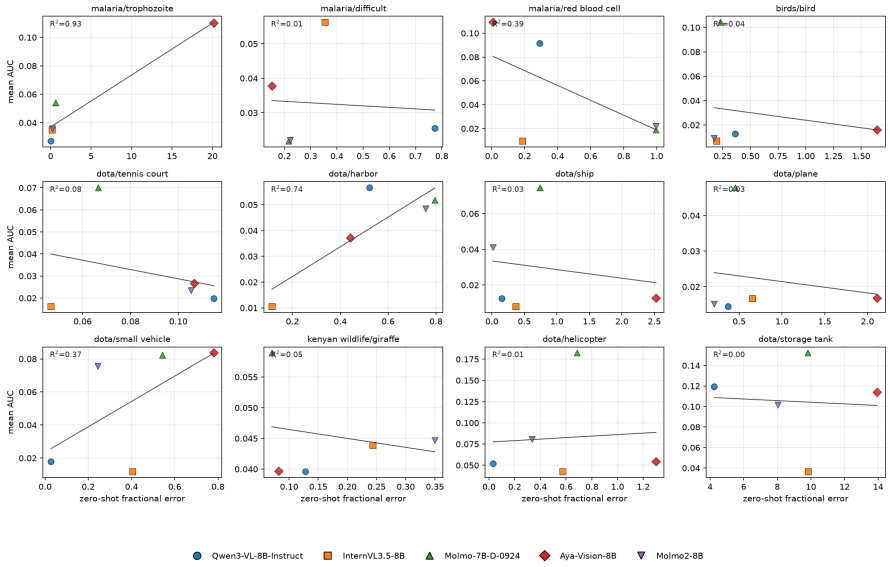

Figure A3: Total-count gap on the full pool vs. mean AUC of DIS (area under the fractional-error curve up to 50% labeled, averaged over 100 trials). Each point is one frozen model on one task; lower is better on both axes.

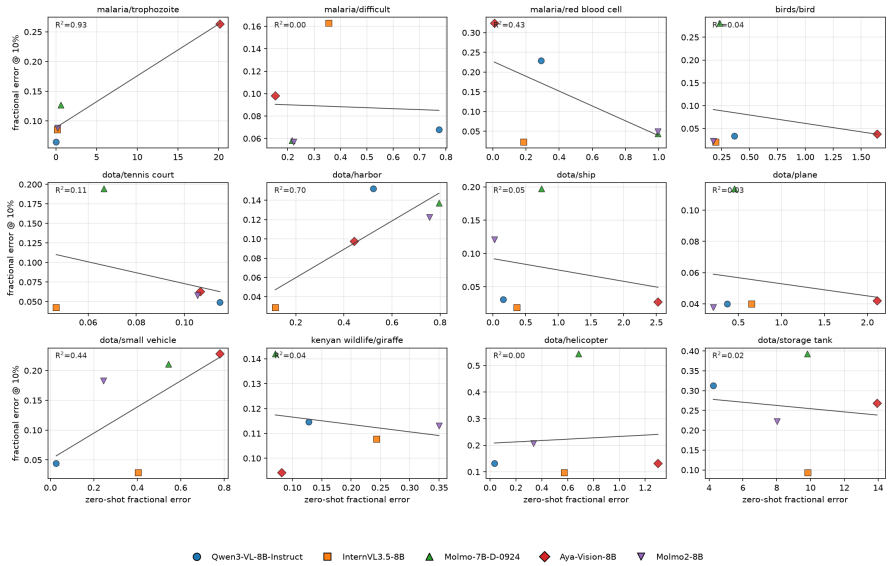

Figure A4: Total-count gap on the full pool vs. fractional error of DIS at 10% labeled. Each point is one frozen model on one task; lower is better on both axes.

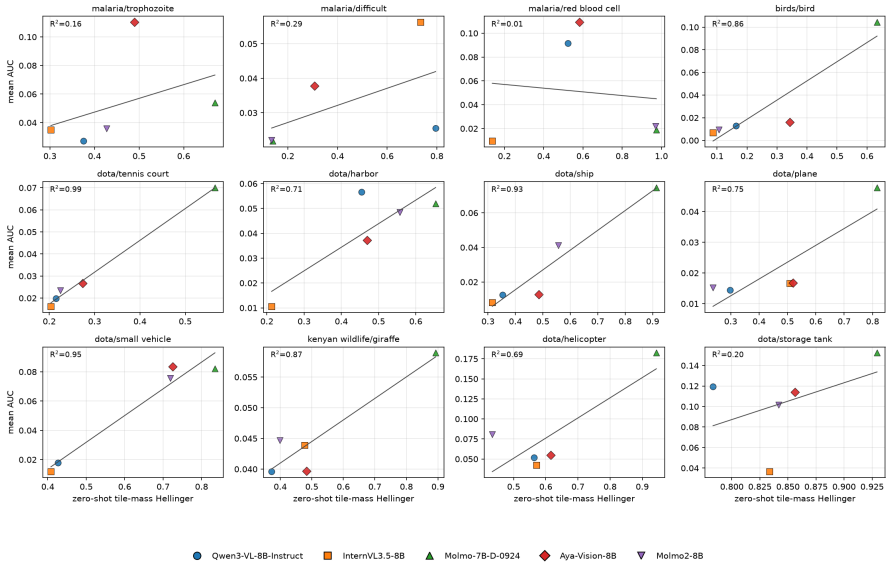

Figure A5: Hellinger distance on the full pool vs. mean AUC of DIS (area under the fractional-error curve up to 50% labeled, averaged over 100 trials). Each point is one frozen model on one task; lower is better on both axes.

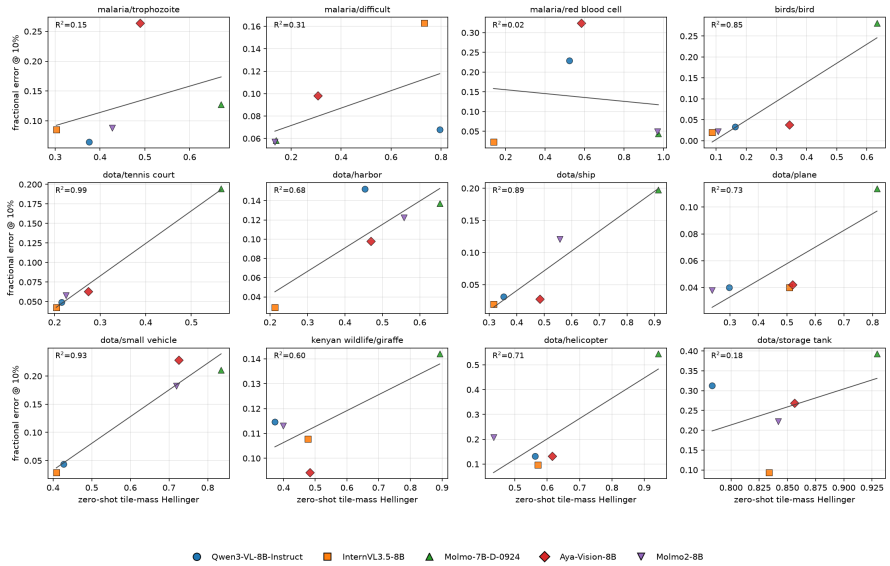

Figure A6: Hellinger distance on the full pool vs. fractional error of DIS at 10% labeled. Each point is one frozen model on one task; lower is better on both axes.

			Hellinger			RMSE			MAE			Count gap		
Candidate set	n	Budget	ρ	beats	ratio	ρ	beats	ratio	ρ	beats	ratio	ρ	beats	ratio
MLLM + OVD	11	10%	**0.70**	**0.80**	**1.06**	0.43	0.73	1.21	0.40	0.70	1.22	0.31	0.56	1.52
		30%	**0.70**	**0.79**	**1.06**	0.45	0.74	1.12	0.39	0.71	1.17	0.25	0.54	1.58
MLLM	8	10%	**0.63**	**0.72**	**1.06**	0.49	0.71	1.15	0.50	0.70	1.15	0.38	0.56	1.32
		30%	**0.62**	**0.71**	**1.04**	0.50	**0.71**	1.06	0.47	0.68	1.12	0.31	0.53	1.53

Table A5: Reliability of four pilot proxies for choosing a frozen model, for two candidate sets of n models at two labeling budgets, over 300 random 50-unit pilots scored by all candidates. ρ is the within-task Spearman correlation between the proxy, averaged over pilots, and the models' realized error, averaged over the 12 tasks. *beats* (mean over pilots) is the fraction of other candidates with higher error than the proxy's pick, and *ratio* ($\times$, median over pilots) is the pick's error divided by the best model's. The ideal value of all three metrics is 1. Bold marks the best proxy for each metric in each row.

Model	$/1,000 units	Units/$1	Normalized
Faster R-CNN			
MobileNetV3	$0.017	59,507	581.3×
ResNet-50	$0.017	57,444	561.1×
Open-vocabulary detectors			
Grounding DINO	$0.126	7,913	77.3×
Florence-2	$0.410	2,440	23.8×
OWLv2	$0.578	1,729	16.9×
Open-source MLLMs			
InternVL3.5-8B	$0.210	4,762	46.5×
Qwen3-VL-8B	$0.238	4,198	41.0×
Molmo-7B-D-0924	$0.240	4,167	40.7×
Molmo2-8B	$0.257	3,894	38.0×
Aya Vision-8B	$0.257	3,885	37.9×
Proprietary MLLMs			
GPT-5.2	$4.645	215	2.1×
Qwen3.7-Plus	$5.912	169	1.7×
Gemini 3.1 Pro	$9.768	102	1.0×

Table A6: Per-unit inference cost, ordered from least to most expensive within each family. *Normalized* is units per dollar relative to Gemini 3.1 Pro. Locally run models assume one A100 at $3 per GPU-hour; proprietary MLLMs use the API cost per unit measured on the 60-unit sample (Table A7).

Model	Origin	Size	GPU-h/unit	API cost/unit	Extraction method
Qwen3-VL-8B	HuggingFace	8B	7.94×10^{-5}	—	Two-pass in-model count→choice
InternVL3.5-8B	HuggingFace	8B	7.00×10^{-5}	—	Two-pass in-model count→choice
Molmo-7B-D-0924	HuggingFace	7B	8.00×10^{-5}	—	Two-pass in-model count→choice
Aya Vision-8B	HuggingFace	8B	8.58×10^{-5}	—	Two-pass in-model count→choice
Molmo2-8B	HuggingFace	8B	8.56×10^{-5}	—	Two-pass in-model count→choice
GPT-5.2	OpenRouter	—	—	$0.004645	"Final answer:" line + GPT-5.2 extractor pass
Gemini 3.1 Pro	OpenRouter	—	—	$0.009768	"Final answer:" line + GPT-5.2 extractor pass
Qwen3.7-Plus	OpenRouter	—	—	$0.005912	"Final answer:" line + GPT-5.2 extractor pass

Table A7: MLLM inference time and cost. Open-source models ran on a single A100 over all 33,605 units of the benchmark, and GPU-hours per unit are seconds per unit divided by 3,600. Proprietary models were queried via OpenRouter on a 60-unit sample, and API cost per unit is the total API cost divided by the number of sampled units.

Prompt	Placeholders and details	Template
Open-source count (stage 1)	The placeholder `{object_name}` is replaced with the task-specific object name (e.g. "birds", "storage tanks").	`How many {object_name} are in this image? Only output the number and no other text.`

Open-source count extraction (stage 2, in-model second pass)	The placeholder `{nlg}` is the free-form text generated in stage 1, and `{max_choice}` is the upper bound of the structured choice list (default 300). The vLLM batched variant omits the "from 0 to {max_choice}" clause.	`What number was this response trying to say?`
		`Response:`
		`{nlg}`
		`Answer only with one integer from 0 to {max_choice}.`

Proprietary count (stage 1)	The placeholder `{object_name}` is replaced with the task-specific object name.	`How many {object_name} are in this image? Think briefly from visible evidence, then finish with a line exactly like Final answer: <number>.`

Proprietary count extraction (stage 2, GPT-5.2 extractor pass)	The placeholder `{response}` is the full reasoning text produced in stage 1.	`Extract the count from the response below.`
		`Response:`
		`{response}`
		`Use the number on the final answer line when present. If the final answer is a range or alternative like '2 or 3', output the first number. Answer only with one number. If the final answer has no digit, answer unknown.`

Verification (optional stage 3, GPT-5.2 verifier pass)	The placeholder `{answer_text}` is the extracted count (or "unknown") and `{response}` is the stage-1 reasoning text.	`Does the response explicitly contain this extracted count as its answer?`
		`Extracted count: {answer_text}`
		`Response:`
		`{response}`
		`Answer yes if the exact count is plainly present in the response. Answer no if the response does not plainly contain that count.`

Table A8: Prompts used for MLLM counting and response processing. Open-source models use a count prompt followed by an in-model count-extraction pass; proprietary models use a count prompt followed by a GPT-5.2 extractor pass and, optionally, a GPT-5.2 verifier pass.

Figure A7: Trajectories of 100 independent AIS (active measurement) runs on the pronghorn survey, each with a different random seed. Every run reached the $\pm 5\%$ estimated-uncertainty stopping condition: 96 after reviewing 3.0–3.4% of the imagery (median 3.2% over all 100 runs), and four after 27.2–28.5%. Of the 100 final 95% confidence intervals, 93 contain the true count; the seven that do not, including all four slow runs, are shown in green.